



Introduction to Docker Orchestration

Course: Containerisation at Scale

Lecture On: Docker Orchestration

Instructor: Dipesh Garg

Prerequisite for this Session

Basic Understanding of Docker

Today's Agenda

- Docker basic concepts - recap
- Setting up an application using Docker Compose
- Issues with Docker Compose
- What and why: Docker orchestration
- Introduction to Docker Swarm and its architecture
- Deploying an application in a Swarm cluster
- Basic features of Swarm
- Docker Compose versus Swarm

What is Docker?

What is Docker?

Rember GTA Vice City Installation?

- Download the installer
- Run the installer
- **Error message during installation**
- Troubleshoot the issue
- Re-run the installer
- **Another issue arises**

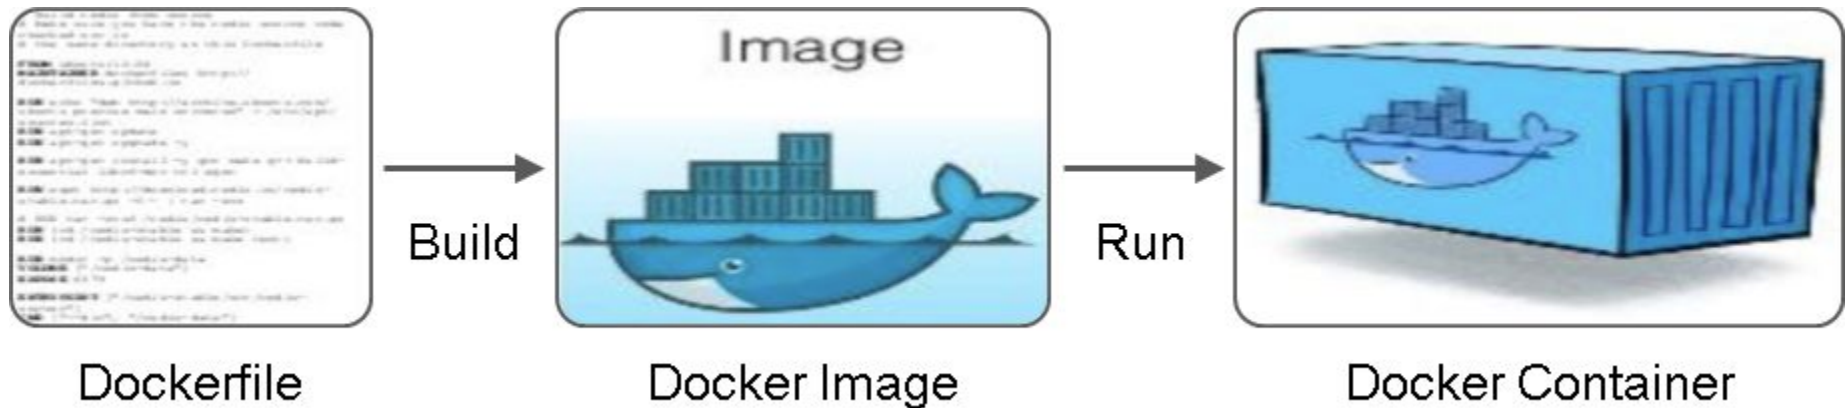
<https://images.app.goo.gl/cSnrBcWHSRwBnmqk9>



- Docker is an open platform. Once we build a Docker container, we can run it anywhere, say, Windows, Linux, a data centre or in a cloud.
- It is a standardised unit, which contains everything that a software application needs to run in any environment, including application code, system libraries and other dependencies.
- It is lightweight, open and secure.

- Virtual Machine
 - It includes multiple applications on a single server.
 - The applications require all the resources and functionalities of the OS.
 - The applications need full isolation and security (since containers provide process-level isolation).
 - Network components that need high performance should be installed on virtual machines (VMs) rather than containers.
- Container
 - It is used to maximise the number of applications running on a server.
 - The applications will be lightweight and the start up time needs to be in milliseconds.
 - The applications need to be infrastructure independent.
 - Containers are more cost efficient than VMs.

- A Dockerfile is a text document that contains all the commands that a user could call in the command line to assemble an image.
- The commands include FROM, RUN, CMD, ADD, COPY, ENTRYPOINT, VOLUME, USER and many more.



Demo 1: Basics of Docker

- Write Dockerfile to create Nginx docker image
- Understanding the use of ADD versus COPY, CMD versus ENTRYPOINT

Poll 1 (15 seconds)

How can you check the list of the space taken by the Docker containers in a system? Like: Images, Containers, Volumes.

- A. `df -h`
- B. `docker df -h`
- C. `docker system df -h`
- D. `docker system df`

Poll 1 (15 seconds)

How can you check the list of the space taken by the Docker containers in a system? Like: Images, Containers, Volumes.

- A. df -h
- B. docker df -h
- C. docker system df -h
- D. docker system df**

Poll 2 (15 seconds)

What is the docker command for displaying the layers of a Docker image?

- A. docker image layers
- B. docker history
- C. docker layers
- D. docker info

Poll 2 (15 seconds)

What is the docker command for displaying layers of a Docker image?

A. docker image layers

B. docker history

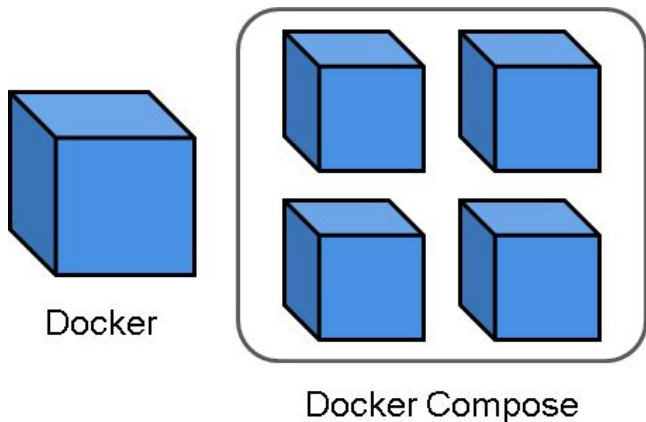
C. docker layers

D. docker info

- Compose is a tool for defining and running multi-container Docker applications.
- With compose, you use a YAML file to configure your application's services.
- You can start all the services with a single command: `docker-compose up`.
- To stop all Docker containers, you can use the command *docker-compose down*.

Here are the features of Docker Compose:

- Multiple isolated environments on a single host
- Preserves volume data when containers are created
- Only recreates containers that have changed




```
version: "3.9" # optional since v1.27.0
services:
  web:
    build: .
    ports:
      - "5000:5000"
    volumes:
      - ./code
      - logvolume01:/var/log
    links:
      - redis
  redis:
    image: redis
volumes:
  logvolume01: {}
```

- They specify the restart policies for how a container should or should not be restarted on exit.
- **Restart Policies**
 - **no**
Do not restart the container (default one) automatically
 - **on-failure**
Restart the container if it exits due to an error
 - **always**
Always restart the container if it stops
 - **unless-stopped**
Always restart the container if it stops, except when done manually

Docker Compose 'Env Variables' & 'Depends On'

- Environment variables can be set in a service's containers with the 'environment' key.

```
web:
  environment:
    - DEBUG=1
    - KEY=VALUE
```

- **Depends On**

- It expresses the dependency between services.
- **docker-compose up** will start the services in the order of dependency. In this example, the 'db' container will begin before the 'web' container:

```
web:
  depends_on:
    - db

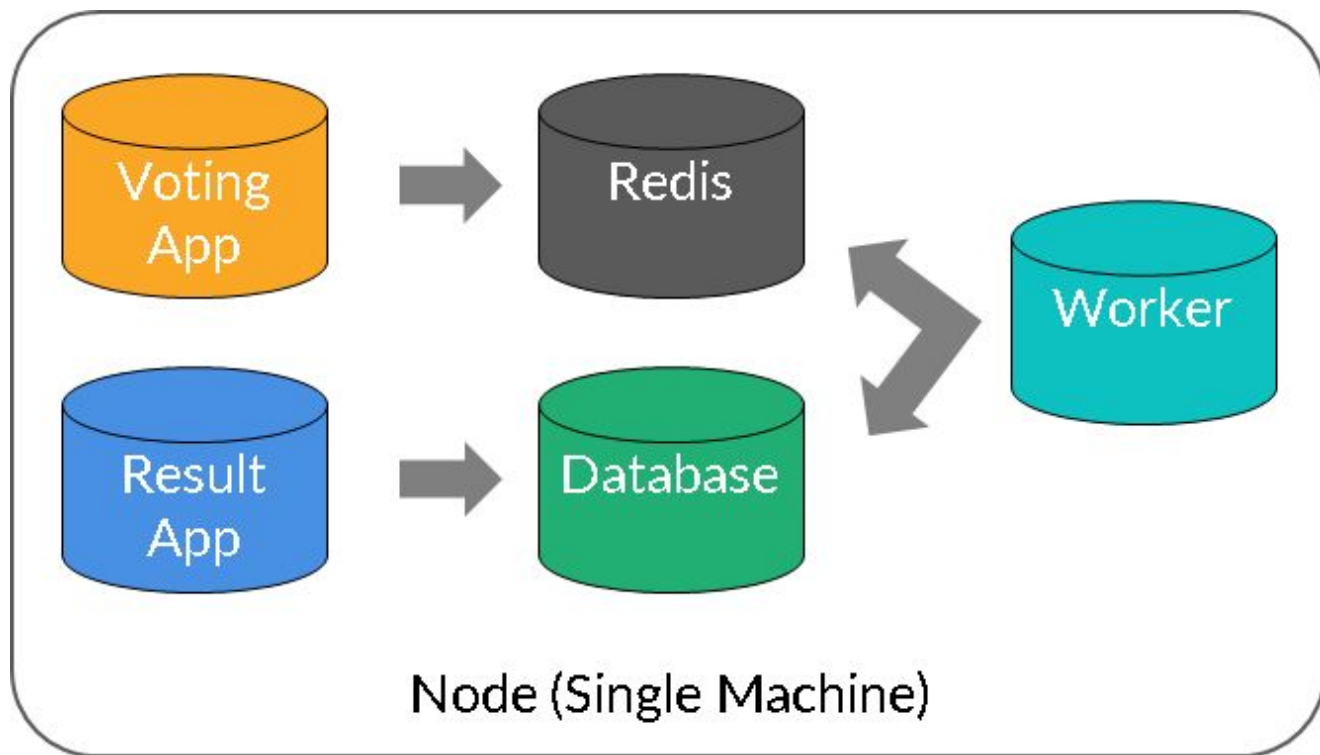
db:
```

- Let's set up a complete application using Docker Compose and see if it solves all of our problems.
- **Voting application:** <https://github.com/docker-samples/example-voting-app>

Demo 2: Docker Compose

- Docker Compose concepts
- Running multiple containers using Docker Compose
- Deploying a voting application:
 - <https://github.com/dockersamples/example-voting-app>
- Issues with Docker Compose

Voting App Using Docker Compose



Poll 3 (15 seconds)

What does the Docker image prune command do?

- A. Displays detailed information on one or more images
- B. Shows the history of an image
- C. Persists a Docker image
- D. Removes unused images

Poll 3 (15 seconds)

What does the Docker image prune command do?

- A. Displays detailed information on one or more images
- B. Shows the history of an image
- C. Persists a docker image
- D. Removes unused images**

Poll 4 (15 seconds)

You want to run an Nginx web server container as a background process. You need to make sure the container's port 80 is mapped to port 8080 on the local host machine. As the final output, you should be able to access the default landing page of the Nginx web server. Which of these is the correct command for this?

- A. `docker run -itd web-server -P 8080:80 nginx`
- B. `docker run -itd web-server -p 80:8080 nginx`
- C. `docker run -itd web-server -P 80:8080 nginx`
- D. `docker run -itd web-server -p 8080:80 nginx`

Poll 4 (15 seconds)

You want to run an Nginx web server container as a background process. You need to make sure the container's port 80 is mapped to port 8080 on the local host machine. As the final output, you should be able to access the default landing page of the Nginx web server. Which of these is the correct command for this?

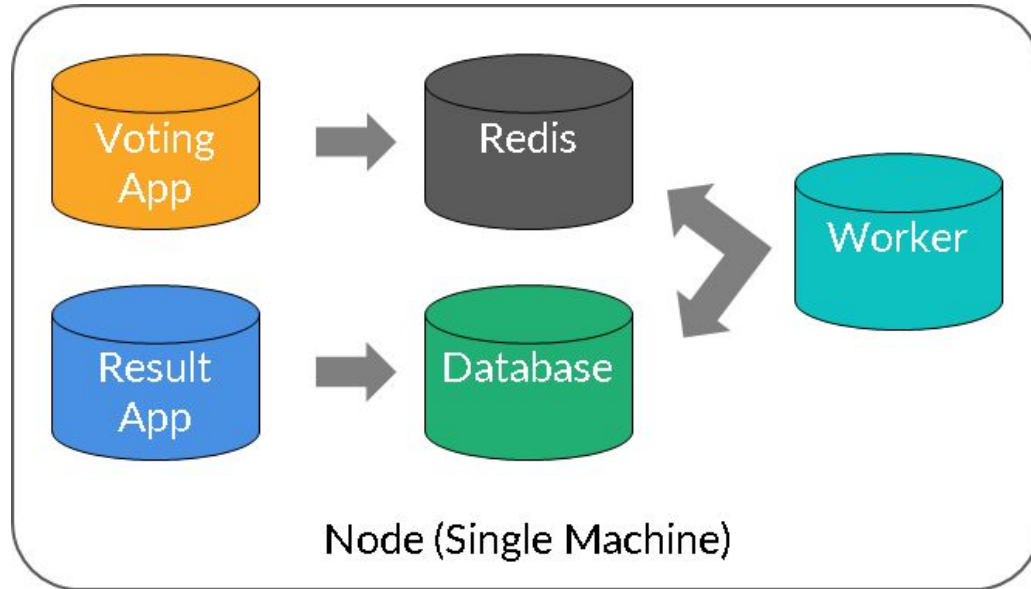
- A. `docker run -itd web-server -P 8080:80 nginx`
- B. `docker run -itd web-server -p 80:8080 nginx`
- C. `docker run -itd web-server -P 80:8080 nginx`
- D. **`docker run -itd web-server -p 8080:80 nginx`**

Container Orchestration

- Container orchestration is all about managing the life cycle of the containers.
- It helps with managing large, dynamic environments.

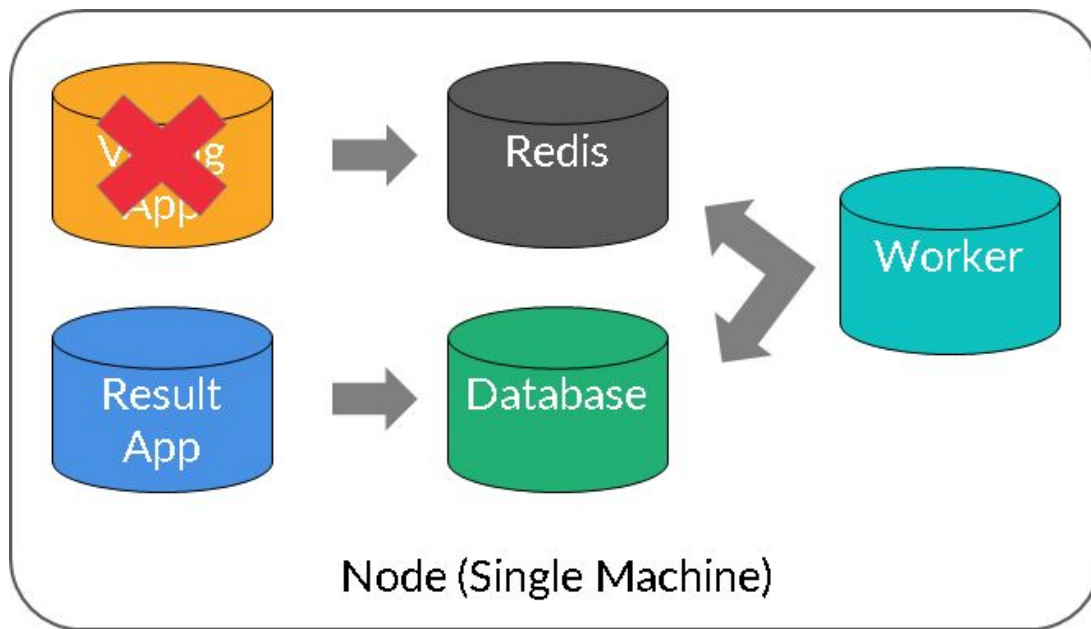


Take the example of the voting application, deployed using Docker Compose on a machine.

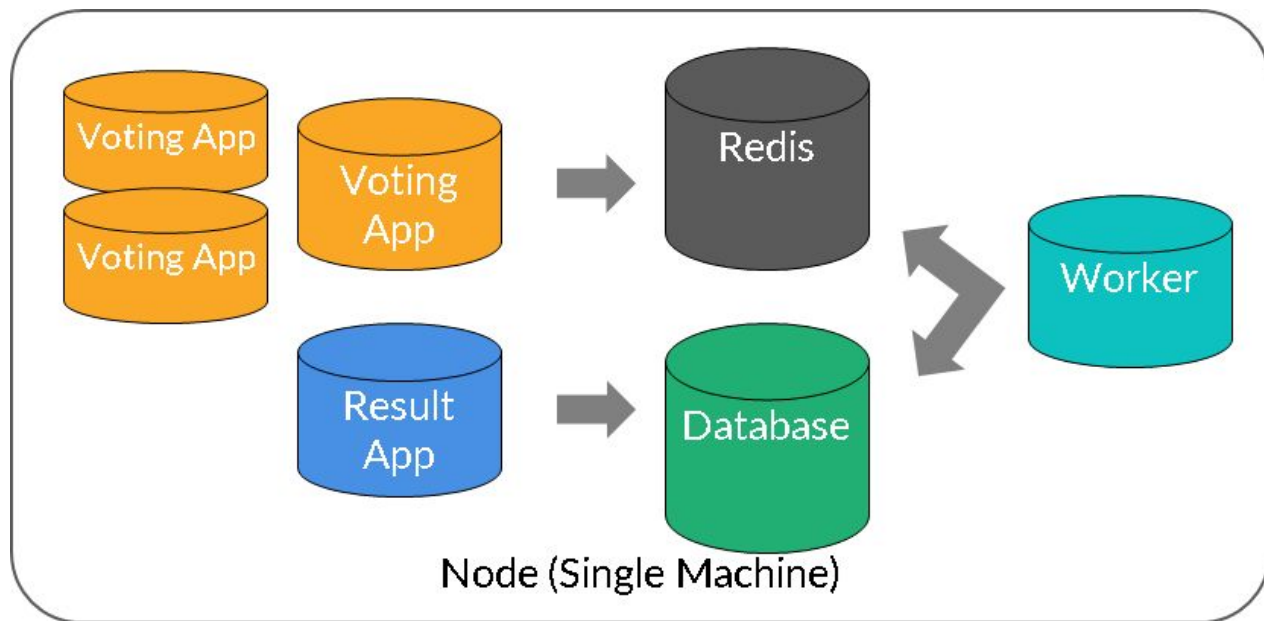


Voting App (Availability Challenge)

What if one of the applications goes down? How we can ensure availability of the application?

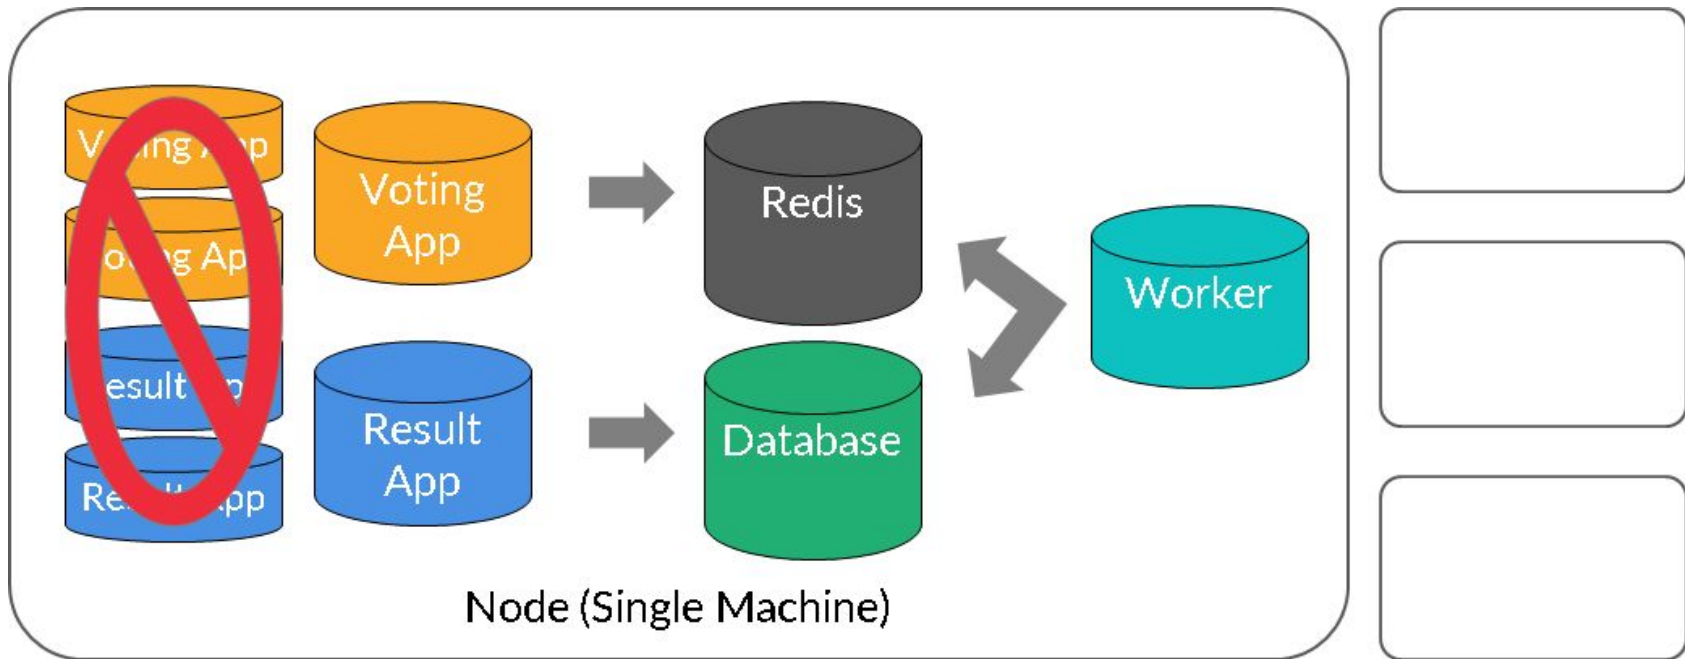


As requests increase, you need to scale your system as well.

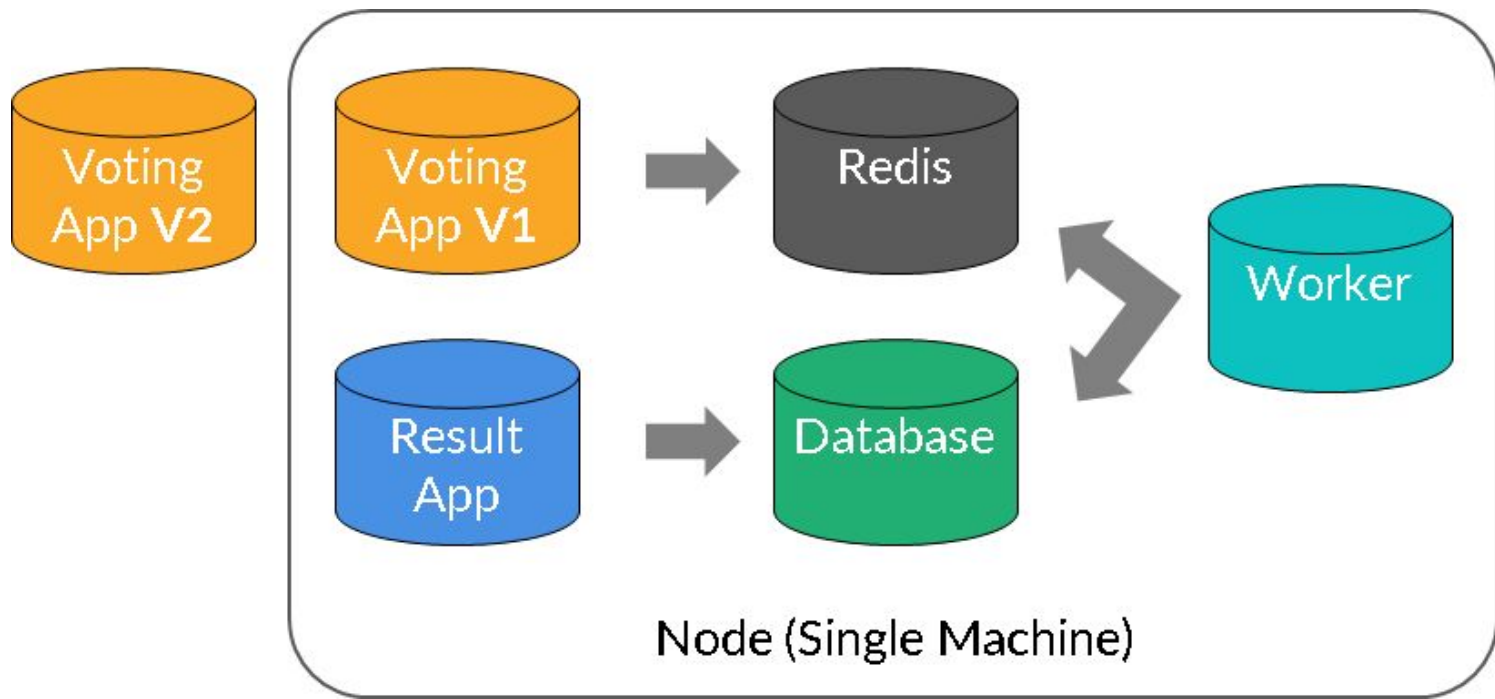


Voting App (Multi-Node Challenge)

One node is insufficient to place a high number of containers.
But how do we shift or deploy containers on other nodes?



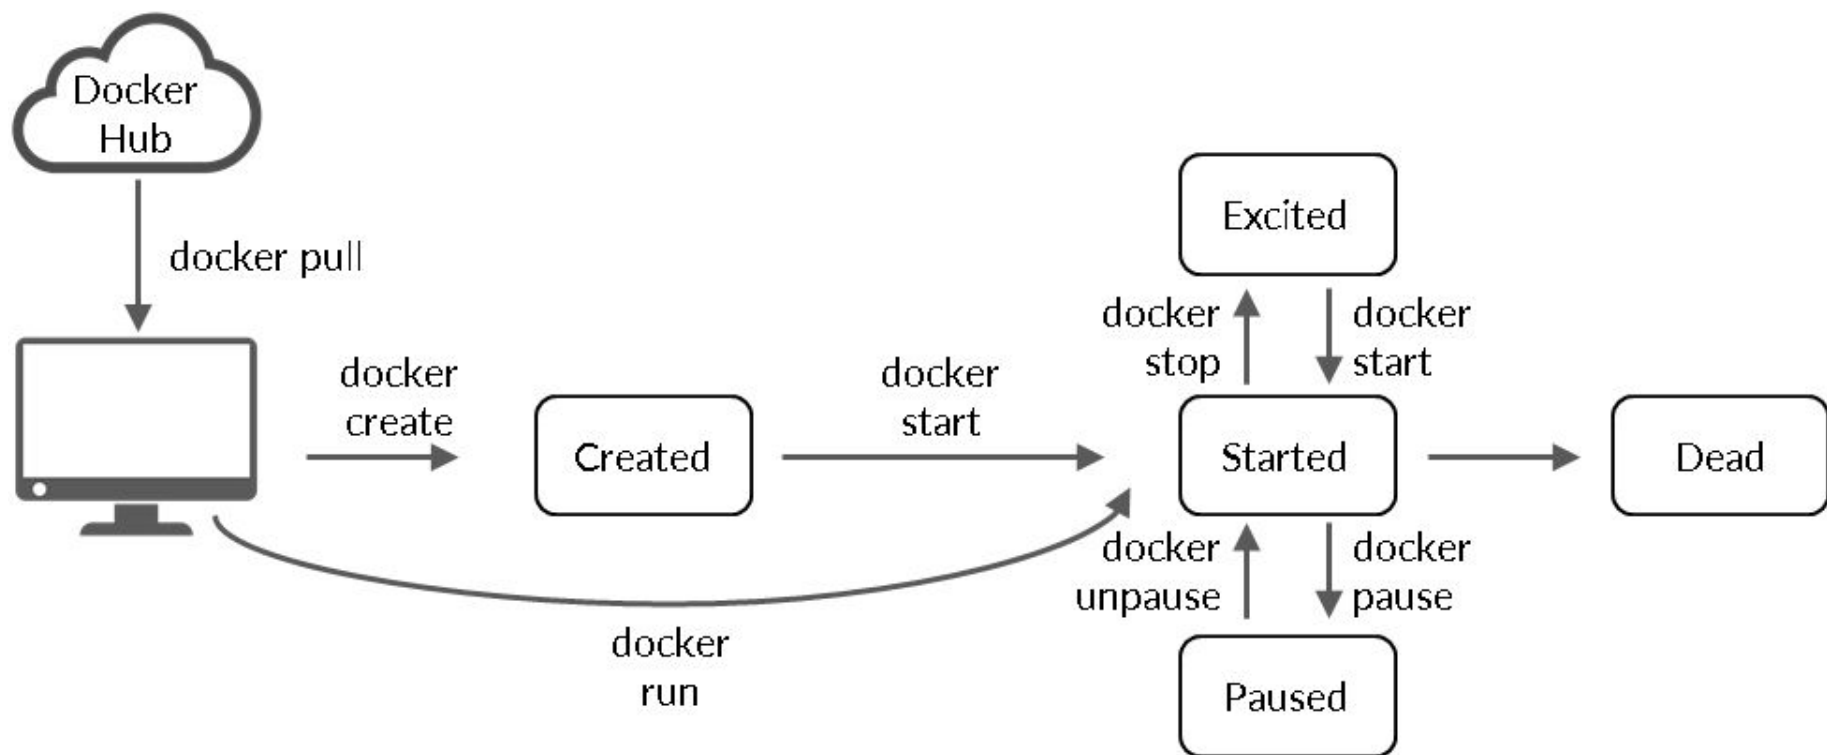
How do we replace V1 with V2 without affecting the end user?



The life cycle of a container includes three phases:

- **Starting phase:** Configuration phase - Container name, image name, port mapping, volume, network configuration
- **Running phase:** Scaling, load balancing, availability, health check
- **Deletion phase:** Clean-up of the resources associated with a container

Life Cycle of a Container



Container Orchestration

Container
Orchestration Engine
Tool

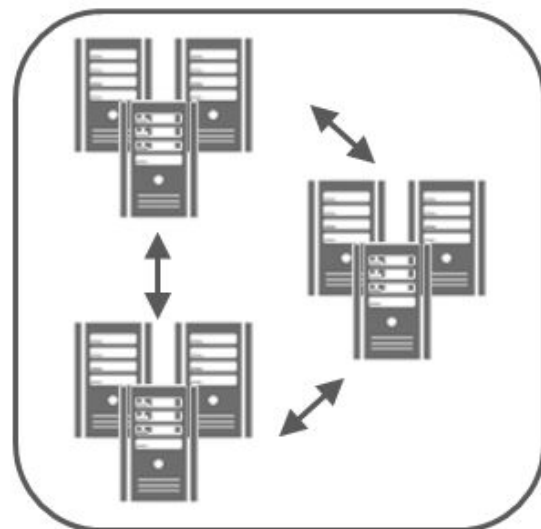


Automate



Configuration
Availability
Load balancing
Scaling
Health check
Security
Networking

Application environment with
multiple servers and
containers



Different Orchestration Tools Available

- Docker Swarm
- AWS managed service - ECS
- K8s - EKS, AKS, GKS
- DigitalOcean Kubernetes Service
- Red Hat OpenShift Online

Docker Swarm

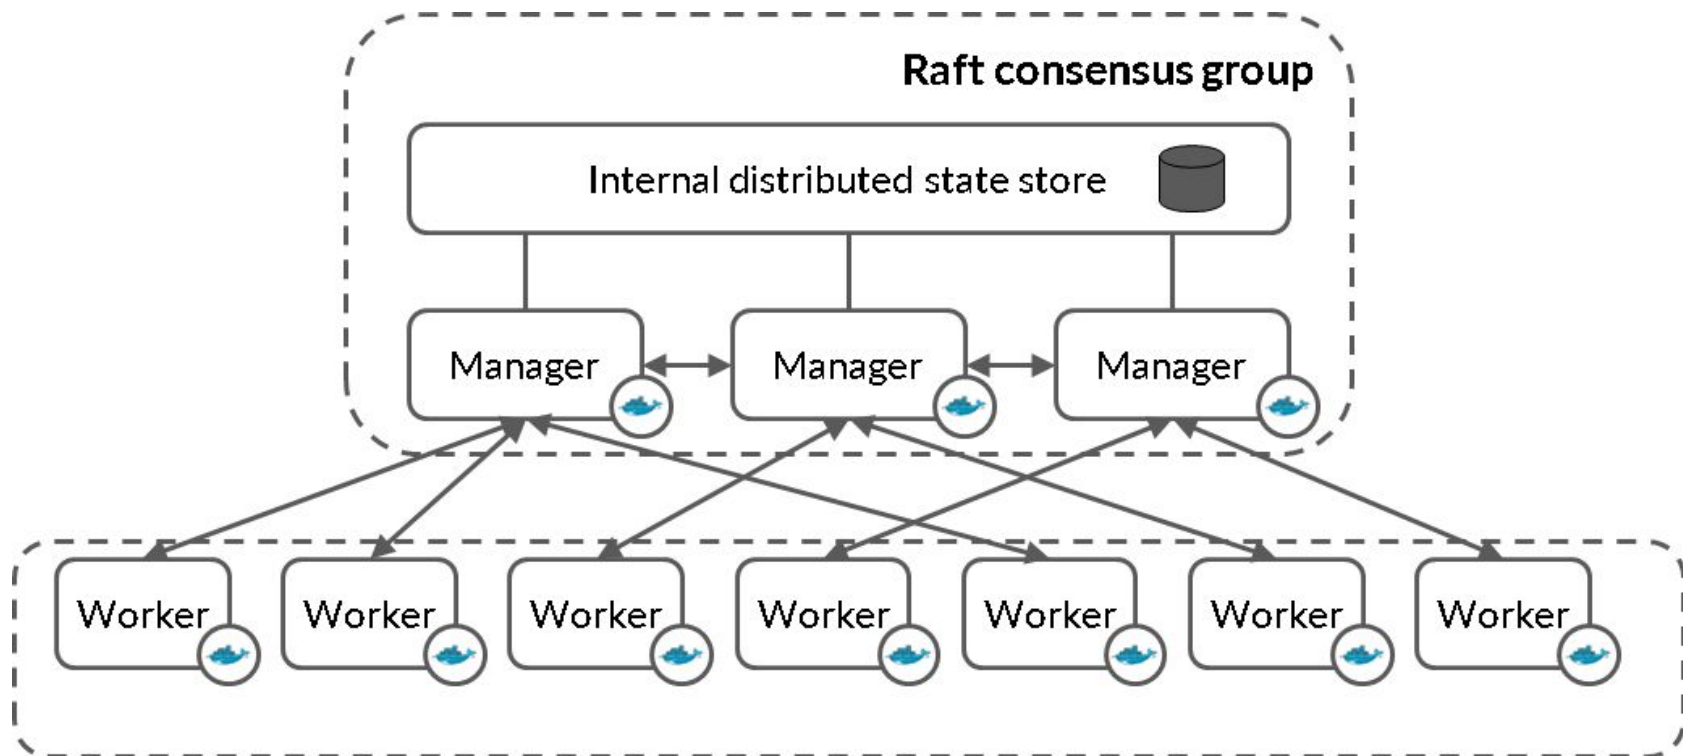
What is Docker Swarm?

Docker Swarm is a container orchestration tool. It is a part of the Docker engine and helps to create, deploy and manage a cluster of Docker containers as a single virtual system.

It helps to manage:

Creation phase -> Running phase -> Deletion phase

- **Nodes** - A node is any machine having a Docker engine installed and participating in a Swarm cluster. We can check all the nodes in a Swarm cluster using the 'docker node ls' command.
- **Manager nodes** - The manager node acts as a manager and assigns units of work called tasks to worker nodes. It also manages cluster-management-level tasks.
- **Leader node** - If you have more than one manager node in the Swarm cluster, then the manager nodes elect a single node called the leader to perform orchestration and schedule tasks.
- **Worker nodes** - These are the slave nodes on which the manager schedules the tasks and the Docker containers run.

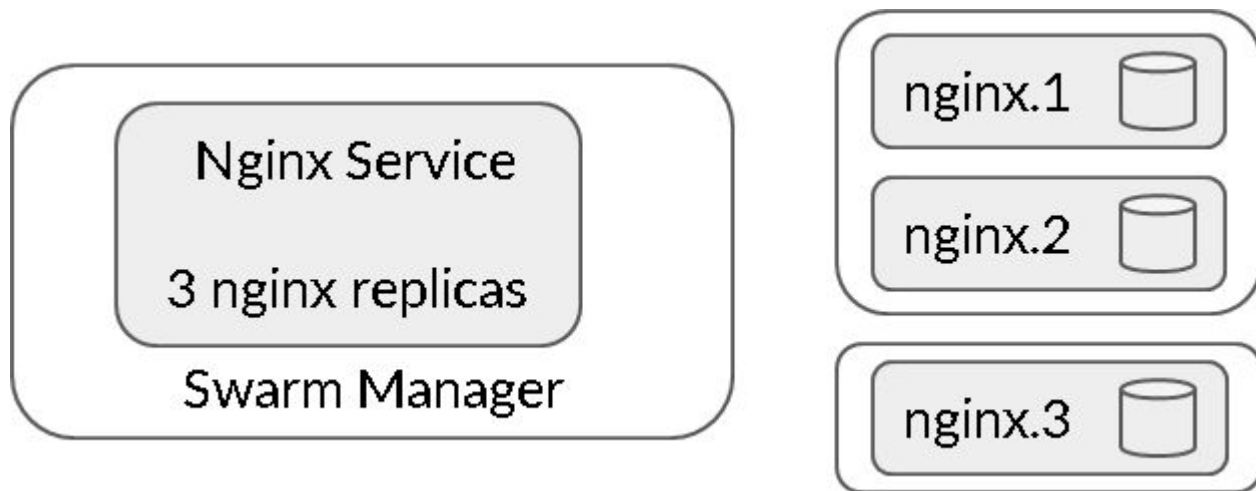


- Docker engine integrated cluster management
- Scaling and load balancing
- Desired state reconciliation
- Multi-host networking
- Provision for security
- Option for rolling updates

- A Swarm cluster consists of master and worker nodes.
- A node is a machine having the Docker engine installed.
- The command for the manager node is:
 - `docker Swarm init --advertise-addr <MANAGER_NODE_IP>`
- The command for the worker nodes is:
 - `docker Swarm join --token <TOKEN>`

- **Task** - Carries the container and describes the container/blueprint:
 - Assume this to be the “Starting phase of the container”
 - Defines the container name, image, port mapping, logging
 - Contains all other container configurations, such as network, logging and CMD command.
- **Service** - Definition of the task to execute on the manager/worker node:
 - Assume this to be the “Running phase of the container”
 - **Defines:** The number of containers to run of a given task
 - Ensures availability, scaling and load balancing
 - Example: Run three containers of Nginx

```
docker service create --name webserver --replicas 3 nginx
```



- Execute these commands in a Swarm cluster and observe the output:
 - `docker node ls`
 - `docker service ls`
 - `docker service create --name webserver --replicas 2 nginx`
 - `docker service ps webserver`
 - `systemctl stop docker` in one of the worker nodes
 - `docker service ps`

```
ubuntu@ip-172-31-5-46:~$ docker service ps webserver
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
u2iyjkqgjac2	webserver.1	nginx:latest	ip-172-31-1-160	Running	Running 8 seconds ago		

```
ubuntu@ip-172-31-5-46:~$
```

```
ubuntu@ip-172-31-5-46:~$
```

```
ubuntu@ip-172-31-5-46:~$
```

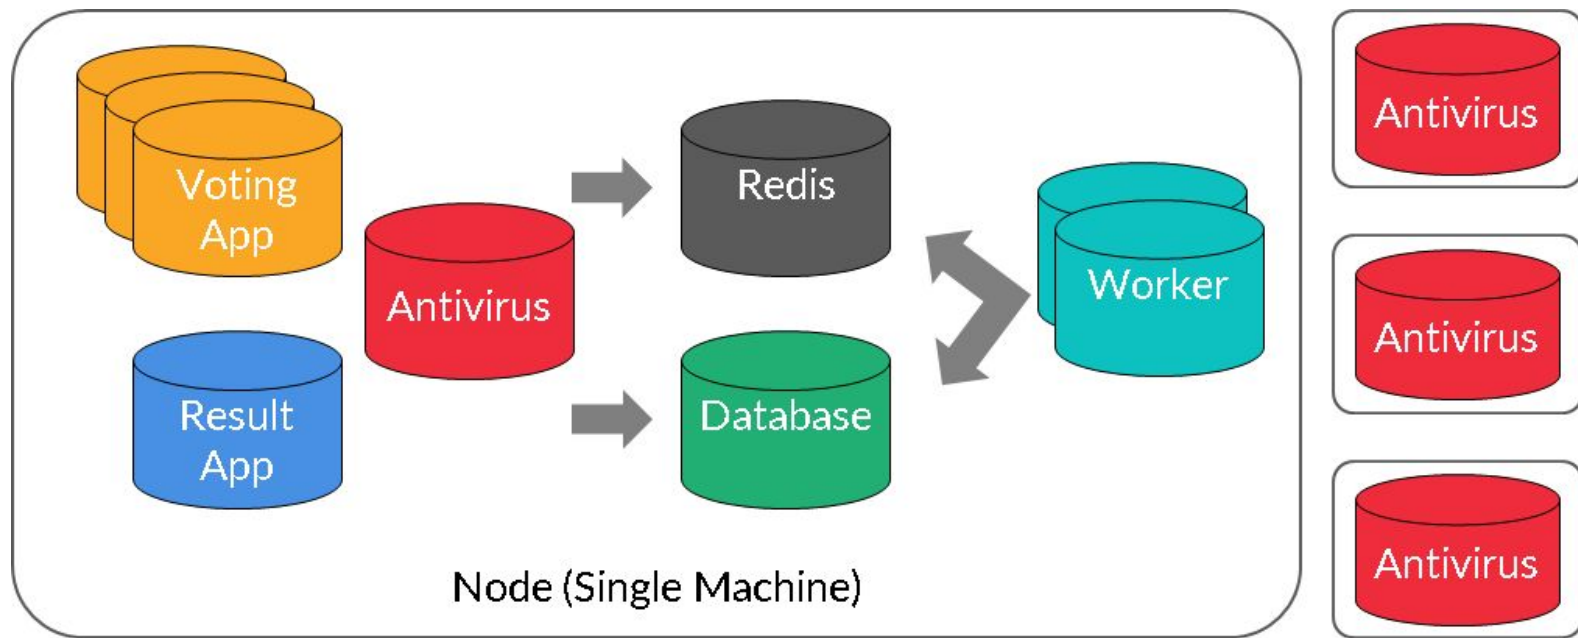
Docker Swarm Service Scaling

- Execute these commands in a Swarm cluster and observe the output:
 - `docker node ls`
 - `docker service ls`
 - `docker service create --name webserver --replicas 2 nginx`
 - `docker service ps webserver`
 - `docker service scale webserver=4`

```
ubuntu@ip-172-31-5-46:~$
ubuntu@ip-172-31-5-46:~$ docker service scale webserver=4
webserver scaled to 4
overall progress: 4 out of 4 tasks
1/4: running [=====]
2/4: running [=====]
3/4: running [=====]
4/4: running [=====]
verify: Service converged
ubuntu@ip-172-31-5-46:~$ docker service ls
ID                NAME        MODE           REPLICAS  IMAGE          PORTS
rws2pzntalqt     webserver   replicated     4/4        nginx:latest
ubuntu@ip-172-31-5-46:~$ docker service ps webserver
ID                NAME        IMAGE          NODE           DESIRED STATE  CURRENT STATE      ERROR          PORTS
u2iyjkqggjac2    webserver.1  nginx:latest   ip-172-31-1-160 Running        Running 4 minutes ago
vklus60ckr5m3    webserver.2  nginx:latest   ip-172-31-5-46  Running        Running 32 seconds ago
way8dqi4nwd5     webserver.3  nginx:latest   ip-172-31-5-46  Running        Running 32 seconds ago
icapf41xlwlp     webserver.4  nginx:latest   ip-172-31-1-160 Running        Running 35 seconds ago
ubuntu@ip-172-31-5-46:~$
```

Replicated Versus Global Service

- **Replicated** - You specify the number of identical tasks that you want to run.
- **Global** - It is a service that runs one task on every node.



Replicated Versus Global Service

- **Replicated**

- docker service create --name replica --replicas 1 nginx
- docker service ps

```
ubuntu@ip-172-31-5-46:~$ docker service ls
ID                NAME      MODE      REPLICAS  IMAGE      PORTS
ebrsyr16pj9l     replica   replicated 1/1        nginx:latest
```

- **Global**

- docker service create --name antivirus --mode global -dt ubuntu
- docker service ps

```
ubuntu@ip-172-31-5-46:~$ docker service ps antivirus
ID                NAME      IMAGE      NODE           DESIRED STATE  CURRENT STATE      ERROR      PORTS
kwjhbto26whp     antivirus.ubuntu:latest ip-172-31-1-160 Running        Running 15 seconds ago
zlk47yvp6xer6     antivirus.ubuntu:latest ip-172-31-5-46  Running        Running 19 seconds ago
```

```
ubuntu@ip-172-31-5-46:~$
ubuntu@ip-172-31-5-46:~$
ubuntu@ip-172-31-5-46:~$ docker service ls
ID                NAME      MODE      REPLICAS  IMAGE      PORTS
na8op9ho2m6a     antivirus  global     2/2        ubuntu:latest
```

Poll 5 (15 seconds)

Jamie works in the security team of his company, and as part of a log monitoring solution, he wants to push all the logs from the server to a central log-monitoring stack. He has created a container, which can do this automatically. The container should be part of all the servers. The DevOps team is using Swarm as an orchestrator. Which of these is the best way to achieve this?

- A. Use *push all* service mode
- B. Create replicas equal to the number of nodes
- C. Use *global* service mode
- D. Use *replicated* service mode

Poll 5 (15 seconds)

Jamie works in the security team of his company, and as part of a log monitoring solution, he wants to push all the logs from the server to a central log-monitoring stack. He has created a container, which can do this automatically. The container should be part of all the servers. The DevOps team is using Swarm as an orchestrator. Which of these is the best way to achieve this?

- A. Use *push all* service mode
- B. Create replicas equal to the number of nodes
- C. Use *global* service mode**
- D. Use *replicated* service mode

Poll 6 (15 seconds)

Which of these commands is used to get the state of the tasks in a Swarm cluster?

- A. `docker service ls`
- B. `docker service ps`
- C. `docker status`
- D. `docker stats`

Poll 6 (15 seconds)

Which of these commands is used to get the state of the tasks in a Swarm cluster?

- A. `docker service ls`
- B. `docker service ps`**
- C. `docker status`
- D. `docker stats`

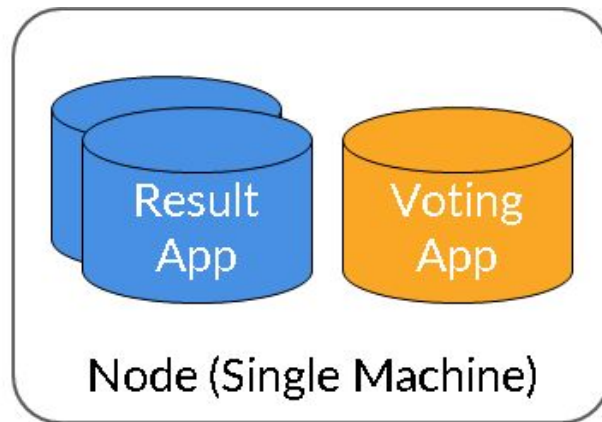
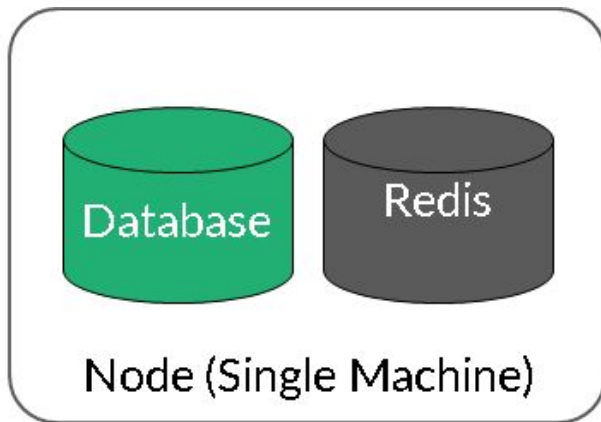
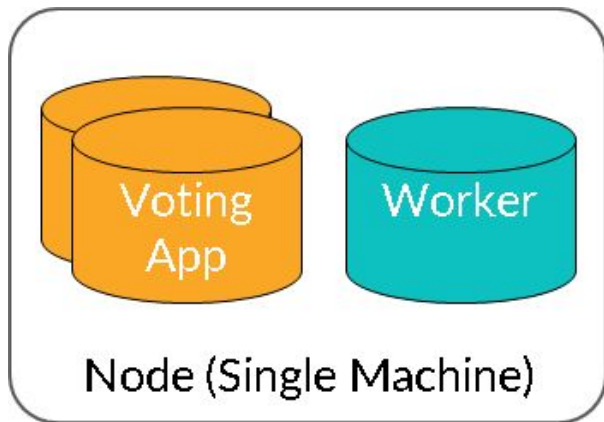
- Replacing a running Docker image version with a new version (replacing V1 with V2)
- During update, ensure that there is no downtime to the application. The end user should not be affected
- Rolling updates in Docker Swarm:
 - Incrementally replace existing containers with containers of a new version of the Docker image
 - **Example:** If four containers are running, first, we can replace two containers with a new version. Once deployed, we can replace the remaining two.

Updates in Docker Swarm

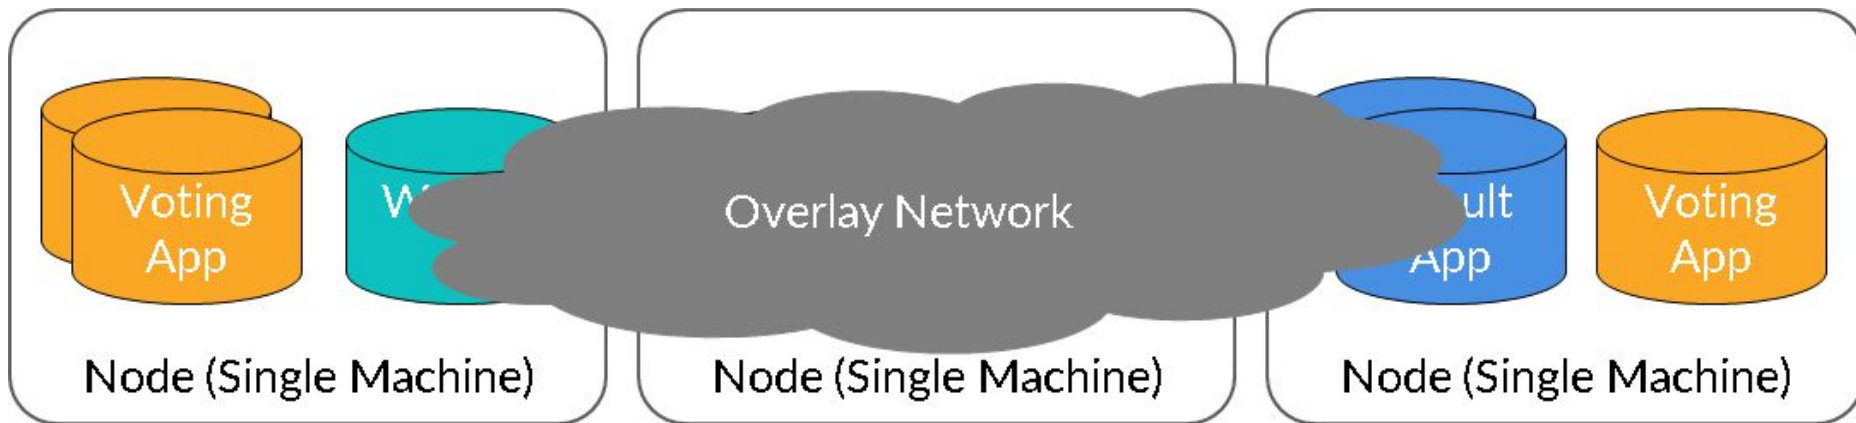
- `docker service create --name webserver --replicas 1 nginx`
- `docker service scale webserver=4`
- `docker service update \`
 - `--update-parallelism 2 \` *(updates two containers at a time)*
 - `--update-delay 20s \` *(configures the time delay between updates)*
 - `--image nginx:1.18-alpine \`
 - `webserver`

```
ubuntu@ip-172-31-5-46:~$  
ubuntu@ip-172-31-5-46:~$ docker service update \  
> --update-parallelism 2 \  
> --update-delay 20s \  
> --image nginx:1.18-alpine \  
> webserver  
webserver  
overall progress: 2 out of 4 tasks  
1/4: running [=====>]  
2/4: running [=====>]  
3/4:  
4/4:
```


- Using the Bridge/Host network, containers on the same host can communicate with each other.
- How will the containers spread across multiple nodes interact?
- You need a new network to communicate with the other containers, which may be spread across any nodes in the Swarm cluster.



- The overlay network driver creates a distributed network among the multiple nodes present in the Swarm cluster.
- It allows the containers connected to it to communicate securely across the cluster.
- When you initialise a Swarm cluster, the overlay network (called ingress) is created by default.



Custom Overlay Network

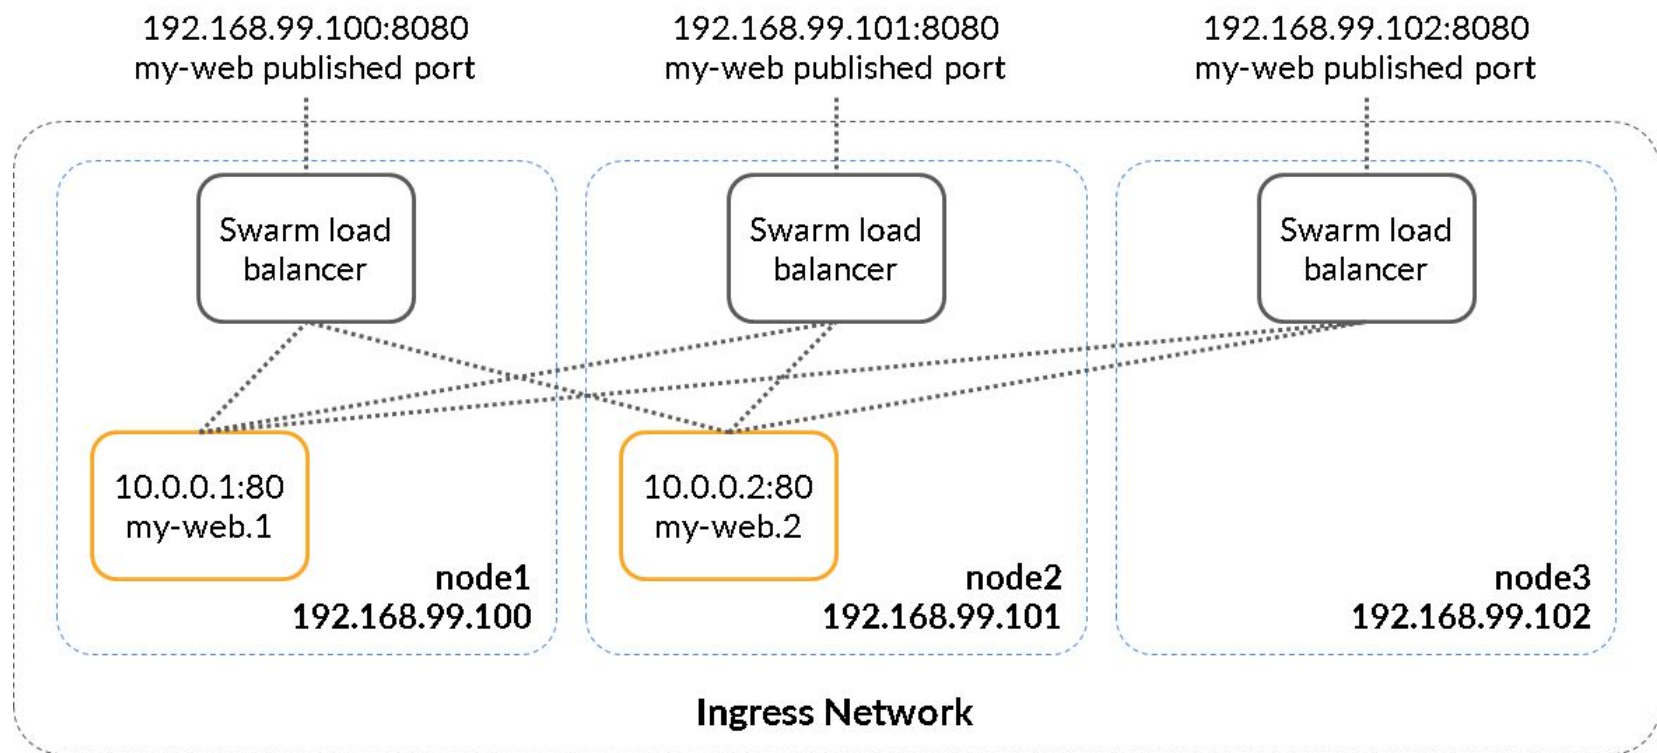
- You can create user-defined overlay networks using docker network create:
 - `docker network create -d overlay my-overlay`
 - `docker network ls`
- You can attach services to the overlay network as shown below:
 - `docker service create --name my-overlay-service --network my-overlay --replicas 3 nginx`

```
ubuntu@ip-172-31-5-46:~$  
ubuntu@ip-172-31-5-46:~$ docker network create -d overlay my-overlay  
9nt9s40dpt4gghnao9m3m03q8  
ubuntu@ip-172-31-5-46:~$ docker network ls  
NETWORK ID          NAME                DRIVER              SCOPE  
6a4aadc36bd9        bridge             bridge              local  
daba1f4a5dff        docker_gwbridge    bridge              local  
6d8bc3414275        host               host                local  
mtc3c1xslwr6        ingress            overlay             swarm  
9nt9s40dpt4g        my-overlay         overlay             swarm  
9ab0271ffbf         none               null                local  
g1t9z3wyrxbg        webserver_default  overlay             swarm  
ubuntu@ip-172-31-5-46:~$  
ubuntu@ip-172-31-5-46:~$
```

- A stack is a group of interrelated services that share the same dependencies and can be orchestrated together.
- Example: Voting application contains Front-End + Back-End + Database.
- A stack can be deployed using a YAML file like the one that we define during docker-compose.
- We can define everything within the YAML file that we might define while creating a docker service.

`docker stack deploy --compose-file docker-compose.yml <NAME>`

Publishing a Port in Swarm



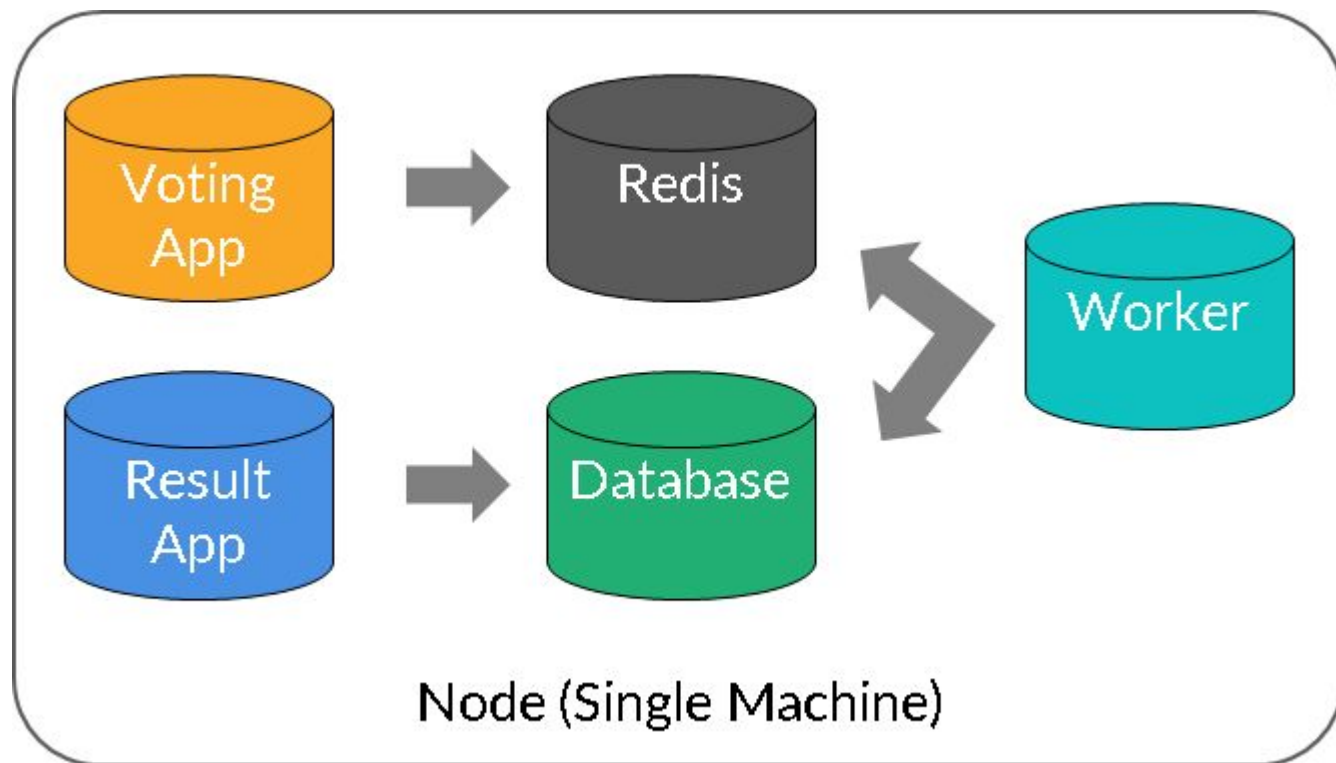
- Let's set up a complete application using Docker Compose and see if it solves all of our problems.
- **Voting application:** <https://github.com/dockersamples/example-voting-app>

```
placement:
  constraints: [node.role == manager]
vote:
  image: dockersamples/examplevotingapp_vote:before
  ports:
    - 5000:80
  networks:
    - frontend
  depends_on:
    - redis
  deploy:
    replicas: 1
    update_config:
      parallelism: 2
    restart_policy:
      condition: on-failure
result:
  image: dockersamples/examplevotingapp_result:before
  ports:
    - 5001:80
```

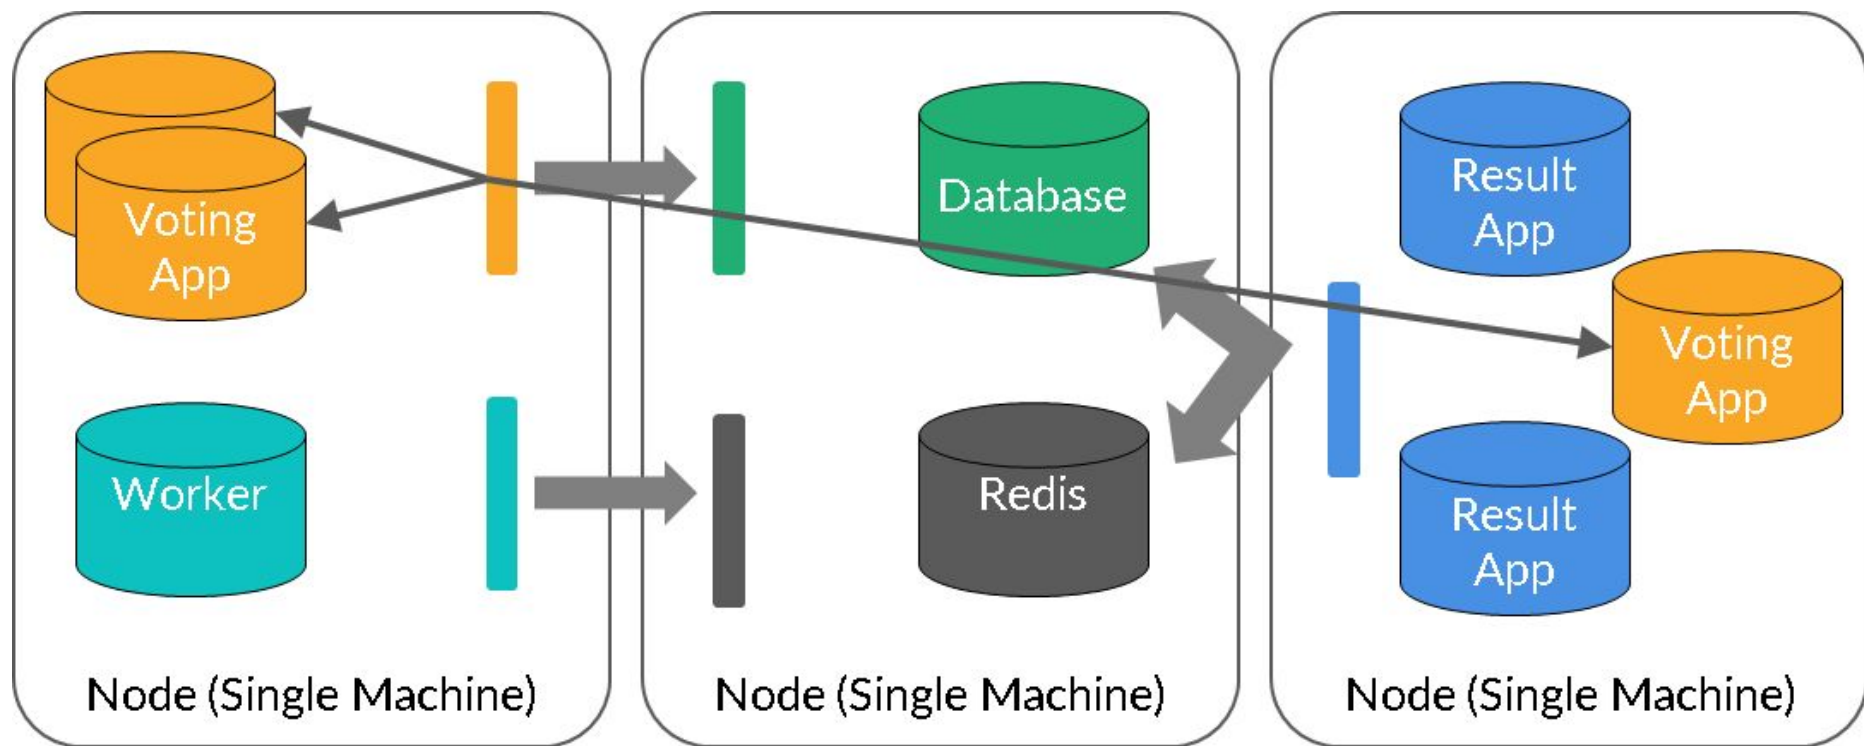
Demo 3: Docker Swarm

- Setting up a three-node Swarm cluster
- Creating, updating and deleting Swarm services
- Overlay network and routing across multiple nodes
- Setting up a voting application using Swarm:
 - <https://github.com/dockersamples/example-voting-app>

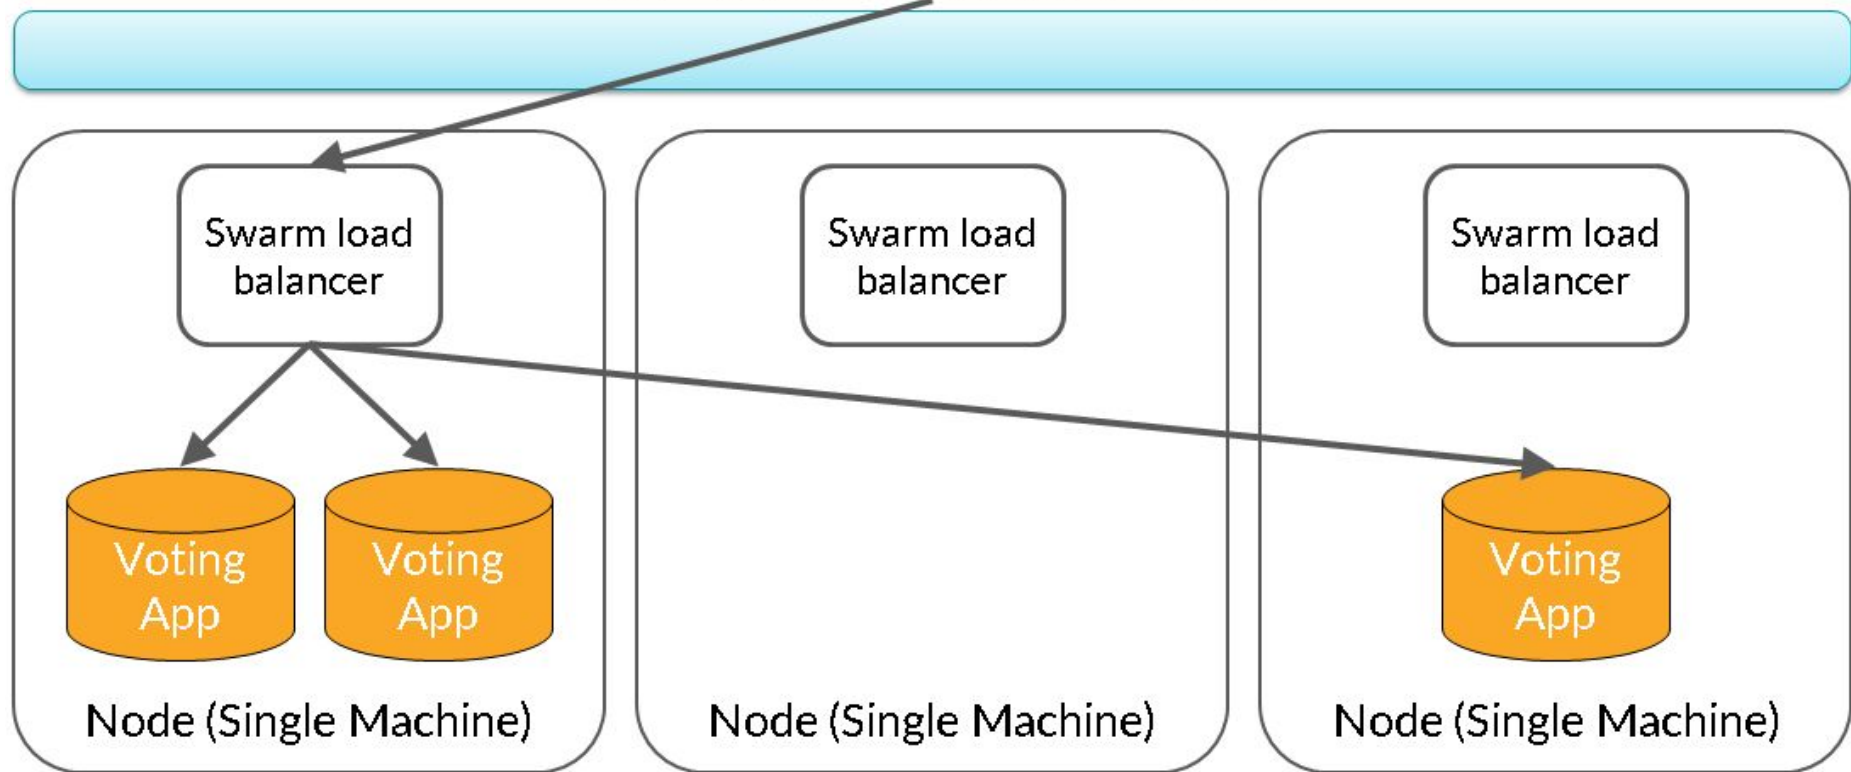
Voting App Using Docker Compose



Voting App Using Docker Swarm



User Hitting Voting App on Node 1



- Swarm provides many features, which include a multi-node setup, scaling, availability, load balancing and many others.
- So, shouldn't we use Swarm always?
- Here are some use cases of Docker Compose:
 - In lower environments like dev/test, you can use Docker Compose.
 - In non-critical applications, you can use Docker Compose to save cost.
 - Docker Compose can be used in urgent and temporary situations.

Important Concepts and Questions

1. What are the differences between 'RUN', 'CMD' and 'ENTRYPOINT' in Dockerfile?
2. What is the Docker build context?
3. In the Docker Compose port mapping definition, Andrew writes 80:80 whereas James writes only 80 for a web service. Can you tell the differences between them?
4. In which situations should you use Docker Compose?
5. What is Docker Orchestration? Name some of orchestration tools available in the market.
6. Can you explain the Swarm architecture and the differences between a service and a task?
7. Jack, a Front-End developer, asked Chris, a DevOps engineer, to scale his application. Chris is confused between the replicated and global service. Can you help him decide?
8. Explain the features of Swarm over running standalone Docker containers.

Doubt Clearance Window

This class has covered the following topics:

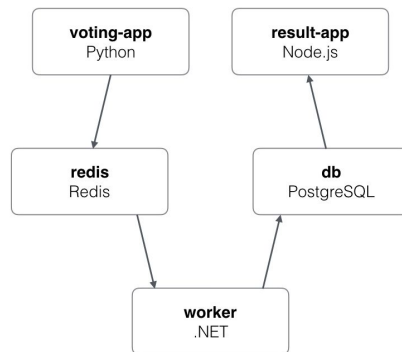
1. Docker basic concepts - recap
2. Setting up an application using Docker Compose
3. Issues with Docker Compose
4. What and why: Docker orchestration
5. Introduction to Docker Swarm and its architecture
6. Deploying an application in a Swarm cluster
7. Basic features of Swarm
8. Docker Compose versus Swarm

What's Next?

1. Nginx: Reverse proxy versus proxy, Nginx folder structure, virtual site hosting
2. Scaling and load balancing
3. Understand why, when and how to scale an application
4. Horizontal scaling versus vertical scaling
5. Load balancing types
6. Schedule services based on different strategies and constraints
7. Limit downtime with rolling updates.
8. Deploy an application on a Swarm cluster

Voting App: <https://github.com/dockersamples/example-voting-app>

1. Set up the voting application using docker-compose. Deploy two replicas of voting frontend using dynamic port mapping. Add an Nginx service (with a customised configuration) file in compose file such that it load balances the request to two voting frontends.
2. Set up the voting application using docker swarm on a two-node cluster by using compose file as a stack file after making required changes. Create separate overlay networks for the frontend and backend. Attach five services of voting application wisely to these networks. Expose the ports of only the frontend services to the host. Set replica count = 5 for voting app and result app



Understand Voting Application

Voting App: <https://github.com/dockeramples/example-voting-app>

1. **Voting-app** : Front-end for casting the vote

Image: dockersamples/examplevotingapp_vote:before

Port to be exposed: 80

2. **Redis**: Db to which voting-app push the votes data

Image: redis:alpine

3. **Worker**: Checks Redis for anything in the queue and push it to PostgreSQL

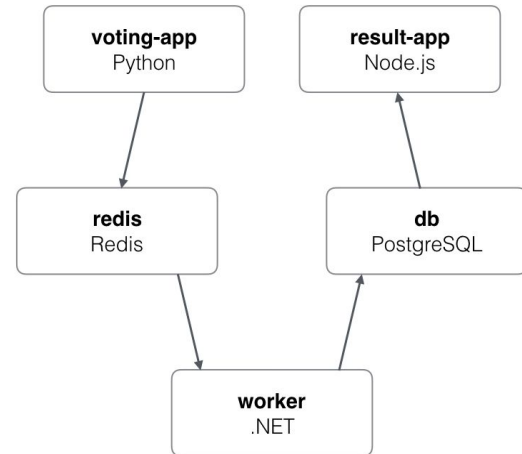
Image: image: dockersamples/examplevotingapp_worker

4. **Db**: Postgresql database

Image: postgres:9.4 ; environment: POSTGRES_USER: "postgres" , POSTGRES_PASSWORD: "postgres"

5. **result-app** : frontend where one can see the result

Image: dockersamples/examplevotingapp_result:before ; Port to be exposed: 80





Thank You!