



Nginx, Scaling and Load Balancing

A Production Architecture Approach

Course: Containerisation at Scale

Lecture On: Nginx and Docker Swarm

Instructor: Dipesh Garg

Prerequisite for this Session

Basic understanding of Docker

Basic understanding of Docker Swarm

Poll 1 (15 seconds)

What is the behaviour expected when a service is created with this command in a swarm cluster containing five nodes?

```
docker service create --publish 8080:80 --replicas=3 nginx
```

- A. Only those nodes which have these containers will listen on port 8080 and will forward the requests to port 80 of any container inside it.
- B. All the nodes will listen on port 80 and will forward the requests to port 8080 of the service.
- C. All the nodes will listen on port 80 and will forward the requests to port 8080 of any one of the containers in the cluster.
- D. All the nodes will listen on port 8080 and will forward the requests to port 80 of any one of the containers in the cluster.

Poll 1 (15 seconds)

What is the behaviour expected when a service is created with this command in a swarm cluster containing five nodes?

```
docker service create --publish 8080:80 --replicas=3 nginx
```

- A. Only those nodes which have these containers will listen on port 8080 and will forward the requests to port 80 of any container inside it.
- B. All the nodes will listen on port 80 and will forward the requests to port 8080 of the service.
- C. All the nodes will listen on port 80 and will forward the requests to port 8080 of any one of the containers in the cluster.
- D. All the nodes will listen on port 8080 and will forward the requests to port 80 of any one of the containers in the cluster.**

Poll 2 (15 seconds)

What would be the output of this PWD instruction of Dockerfile?

```
FROM busybox
```

```
WORKDIR /a
```

```
WORKDIR /b
```

```
WORKDIR c
```

```
RUN PWD
```

- A. /a
- B. /a/b/c
- C. /b/c
- D. /a/bc

Poll 2 (15 seconds)

What would be the output of this PWD instruction of Dockerfile?

```
FROM busybox
```

```
WORKDIR /a
```

```
WORKDIR /b
```

```
WORKDIR c
```

```
RUN PWD
```

A. /a

B. /a/b/c

C. /b/c

D. /a/bc

Today's Agenda

- Nginx: Reverse proxy versus proxy, Nginx folder structure, virtual site hosting
- Scaling and load balancing
- Understand why, when and how to scale an application
- Horizontal scaling versus vertical scaling
- Load balancing types
- Schedule services based on different strategies and constraints
- Deploy an application on a Swarm cluster

Nginx



What is Proxy?

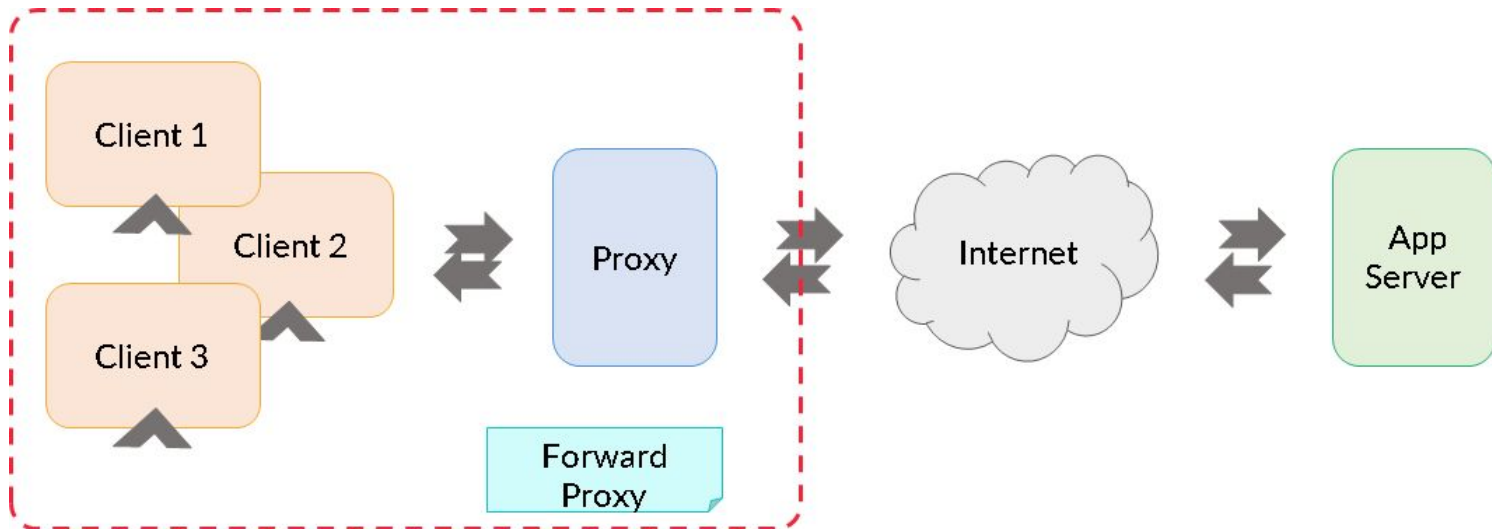
1. In computer networking, a proxy server is a server application or appliance that acts as an intermediary for requests from clients seeking resources from servers that provide those resources.
2. A proxy server thus functions on behalf of the client when requesting service, potentially masking the true origin of the request to the resource server.

Types of Proxy

1. Forward proxy (or just proxy)
2. Reverse proxy

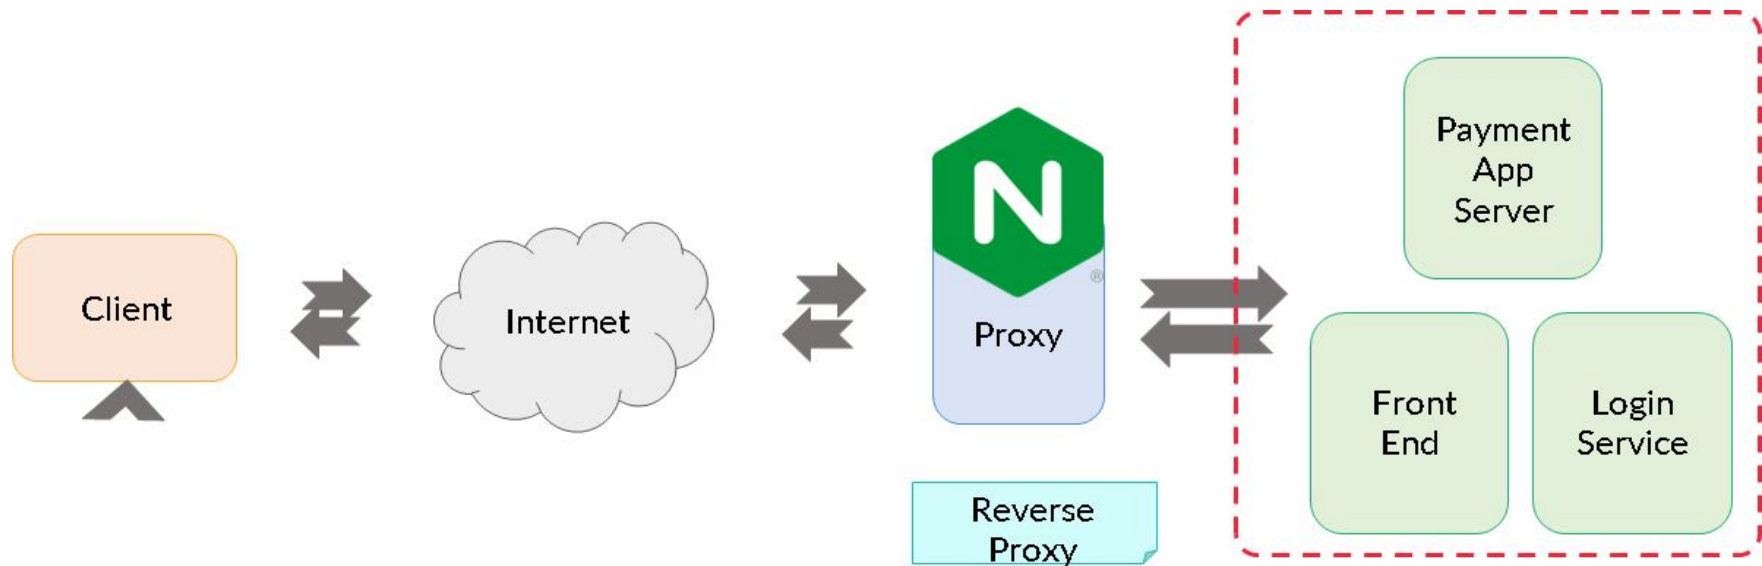
Proxy : Forward Proxy

- A forward proxy provides proxy services to a client or a group of clients. Oftentimes, these clients belong to a common internal network
- This server reveals its identity as a proxy server, but does not disclose the originating IP address of the client to the resource server.



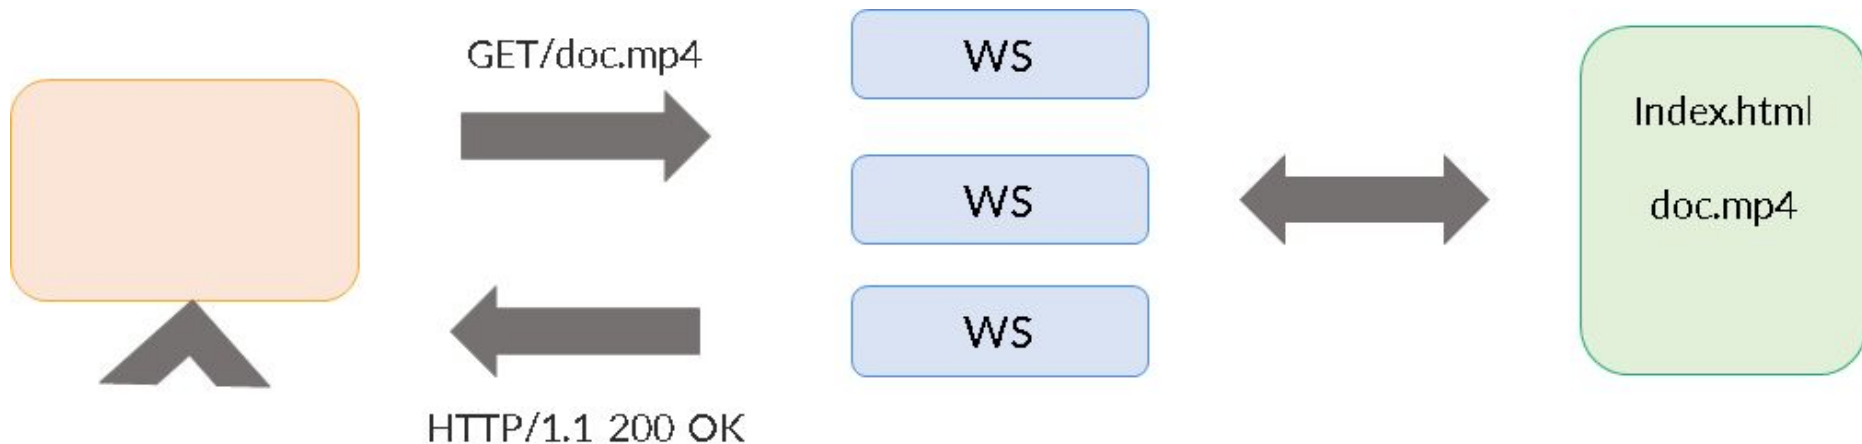
Proxy : Reverse Proxy

- A reverse proxy is a proxy server that appears to clients to be an ordinary server. Reverse proxies forward requests to one or more ordinary servers that handle the request.

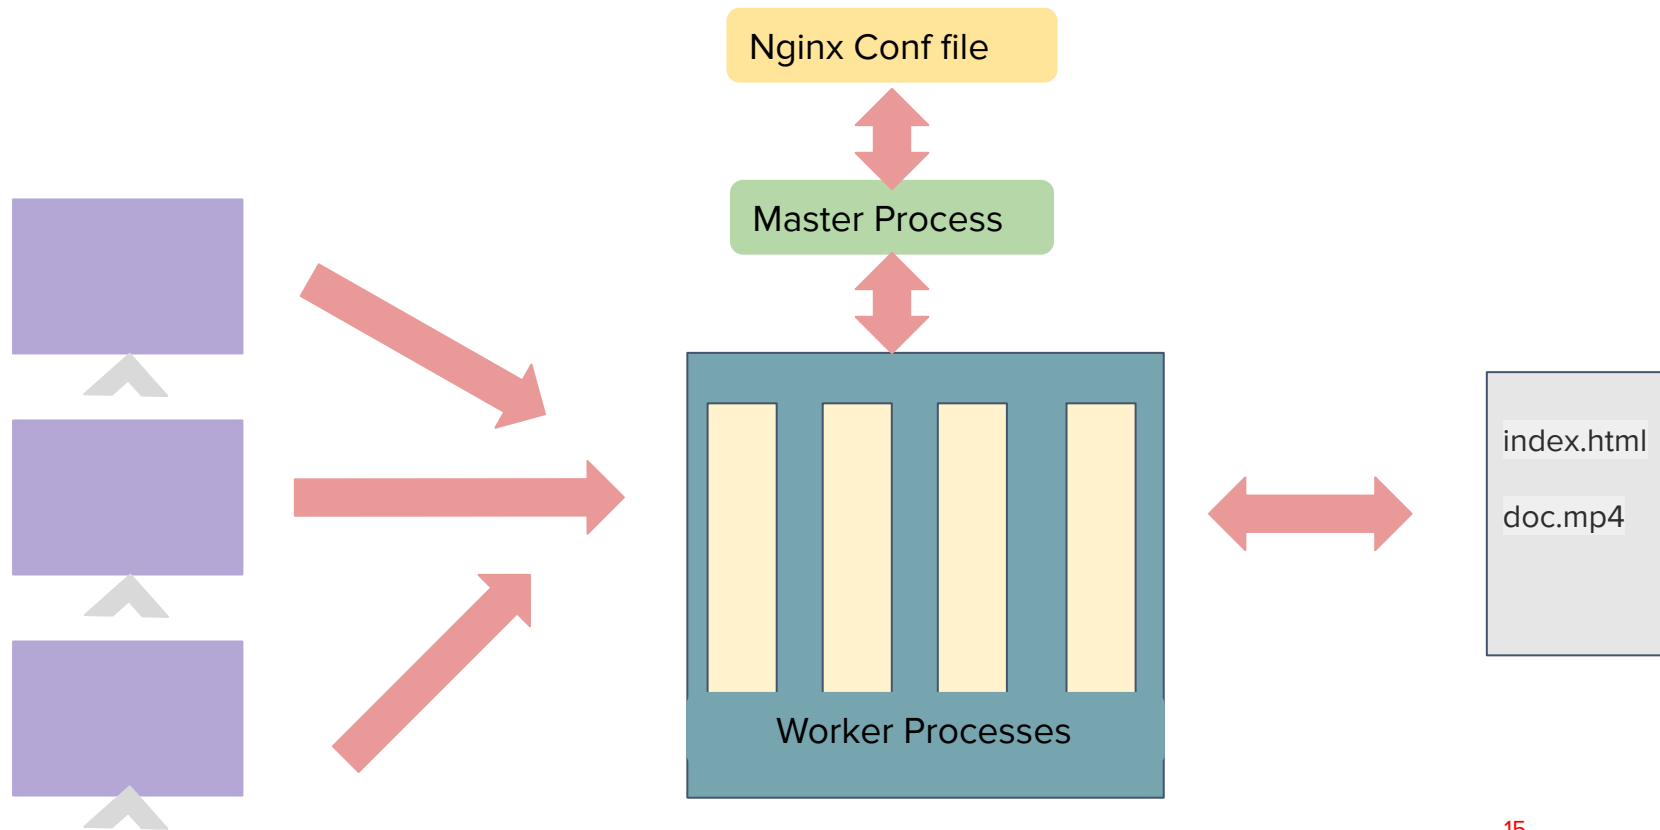


- Forward proxy
 - Anonymises the client's IP
 - Prevents data loss
 - Blocks user from visiting certain websites
 - Improves the user experience by caching external site content
- Reverse proxy
 - Hide the existence of the original back-end servers
 - Load balancing
 - Protect app servers from web based DOS attacks
 - Can provide caching functionality
 - Can optimise content by compressing
 - Combine different websites into a single URL space

- Nginx is one of the most popular web servers in the industry.
- A web server is a program that uses HTTP (Hypertext Transfer Protocol) to serve users the files that form web pages, in response to their requests.



Nginx Architecture



DEMO 1 : NGINX

- Let's take a look at the Nginx folder structure
- Server blocks and custom conf files
- **Host two websites:** Hello upGrad Learners and DevOps is Super Cool
- Serve these websites using upgradlearner.com and devopslearner.com


Nginx configuration file

```
user          www www;
worker_processes 5;  ## Default: 1
error_log     logs/error.log;
pid           logs/nginx.pid;
worker_rlimit_nofile 8192;

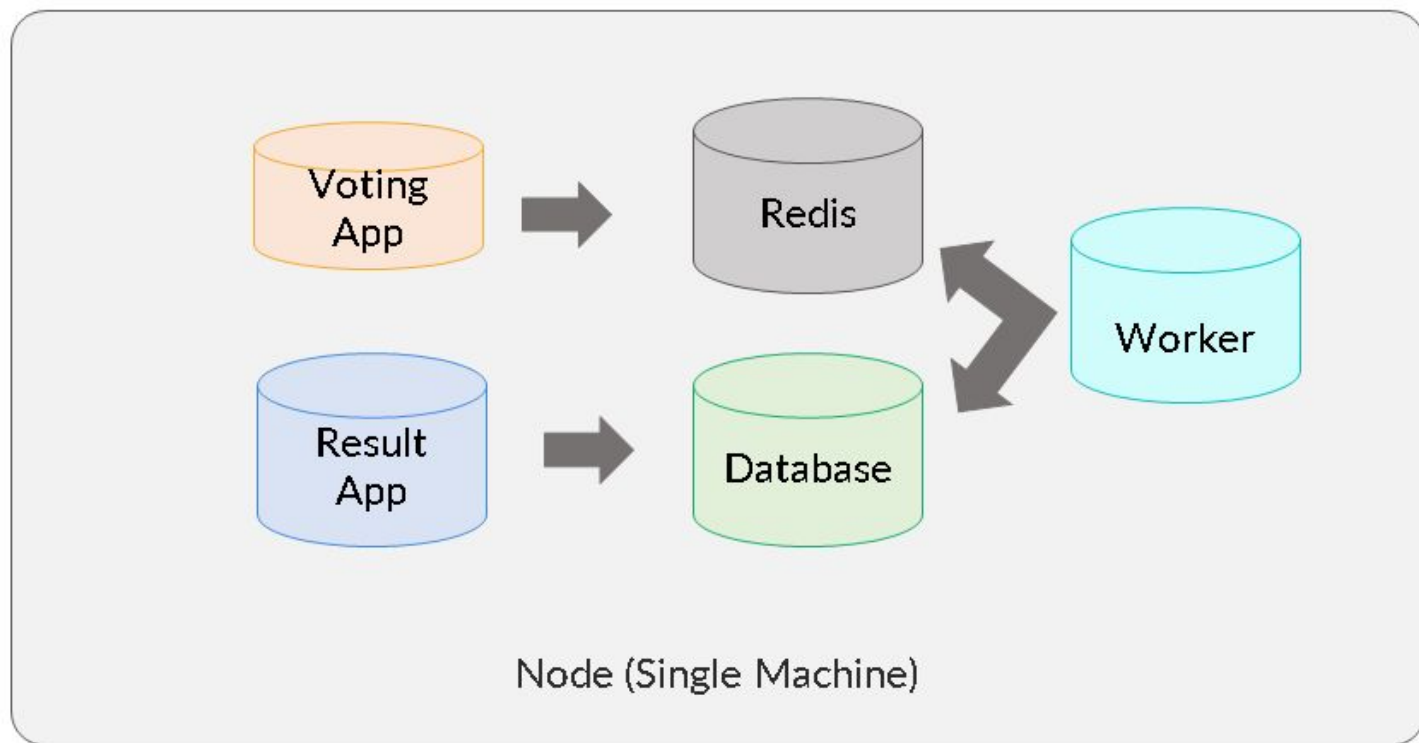
events {
    worker_connections 4096;  ## Default: 1024
}

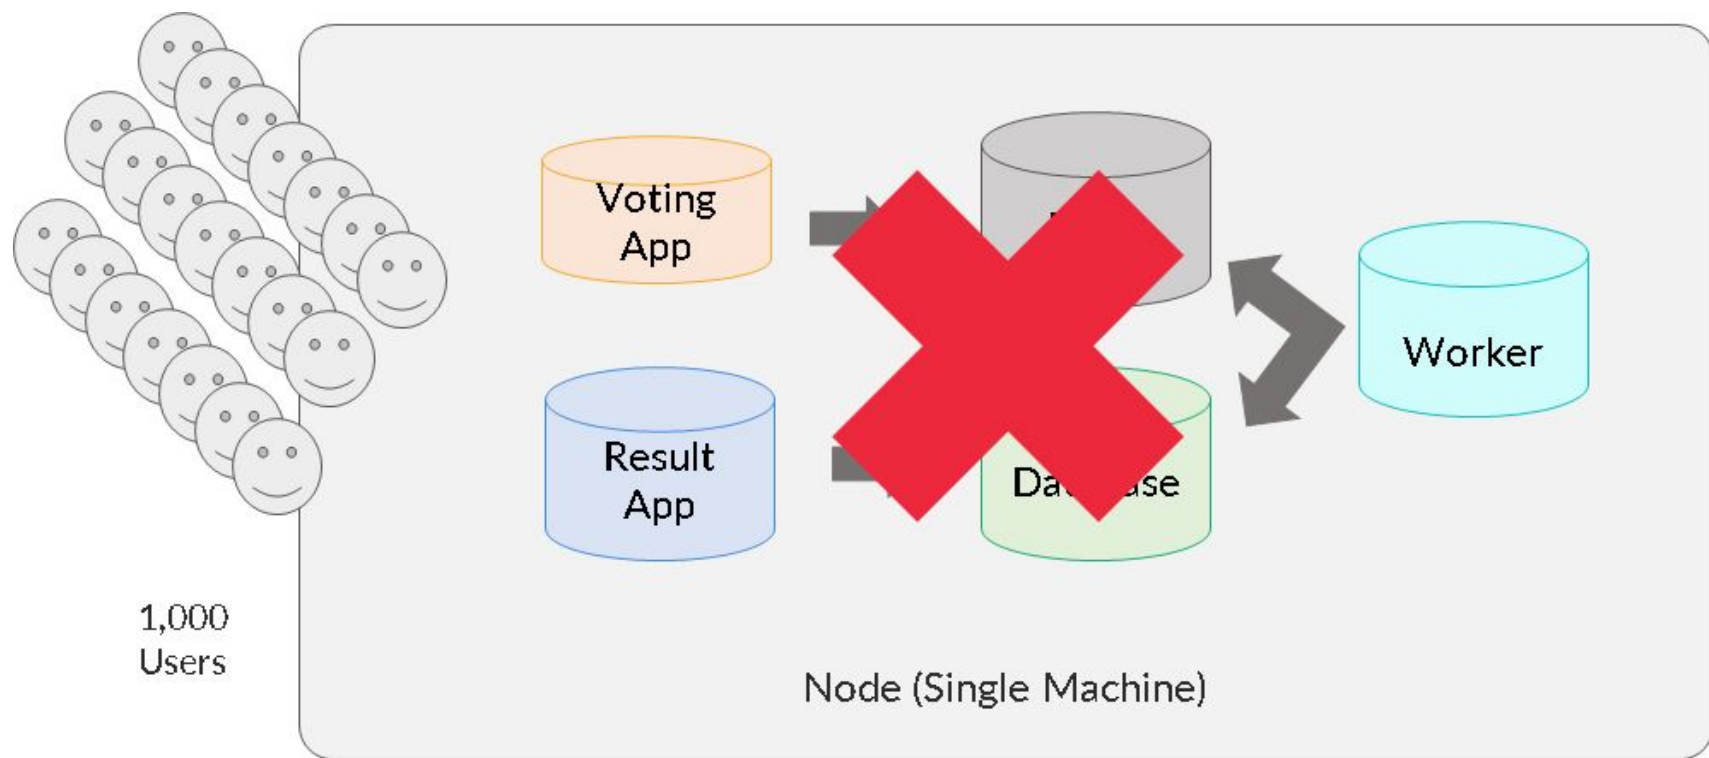
http {
    server { # php/fastcgi
        listen      80;
        server_name  domain1.com www.domain1.com;
        access_log   logs/domain1.access.log  main;
        root         html;
        location ~ /\.php$ {
        }
    }
}
```

Scaling and Load Balancing



100
Users





- Scalability is the ability to easily increase or decrease compute or storage resources as needed to meet changing demand in cost-effective way.
- Ideally if your app takes t seconds to respond to a user request. It should take the same t seconds to respond to each of the million concurrent user requests on your app
- Cloud has simplified all these scaling problems dramatically.
- Remember, scaling is important for all parts of the application. Ignoring any part can result in downtime.

- Vertical scaling
 - Adding resources in existing servers
 - Replacing a server with a bigger/more powerful server
 - Example: Changing t2.small to t2.large in AWS
- Horizontal scaling
 - Provision of additional servers to meet the need
 - Execute concurrent workload and handle traffic efficiently
 - Example: Adding 2 more nodes in a cluster of 3 nodes

Scale Up



Scale Out



Vertical Versus Horizontal

- Vertical
 - When data is on a single machine and partitioning is complicated.
 - The resources perform some cron job operations and duplicacy can lead to errors.
- Horizontal
 - It can provide instant and continuous availability.
 - There is no limit to hardware capacity.
 - Applicable when having multiple machines does not cause data duplicacy errors.
 - You want to run only the required number of hardware and choose not to pay for full capacity every time.

○ Vertical

- Less administrative efforts as you need to manage just one system.
- This involves a downtime.
- Your hardware always runs with full capacity.
- Finite scope of upgradeability in the future.

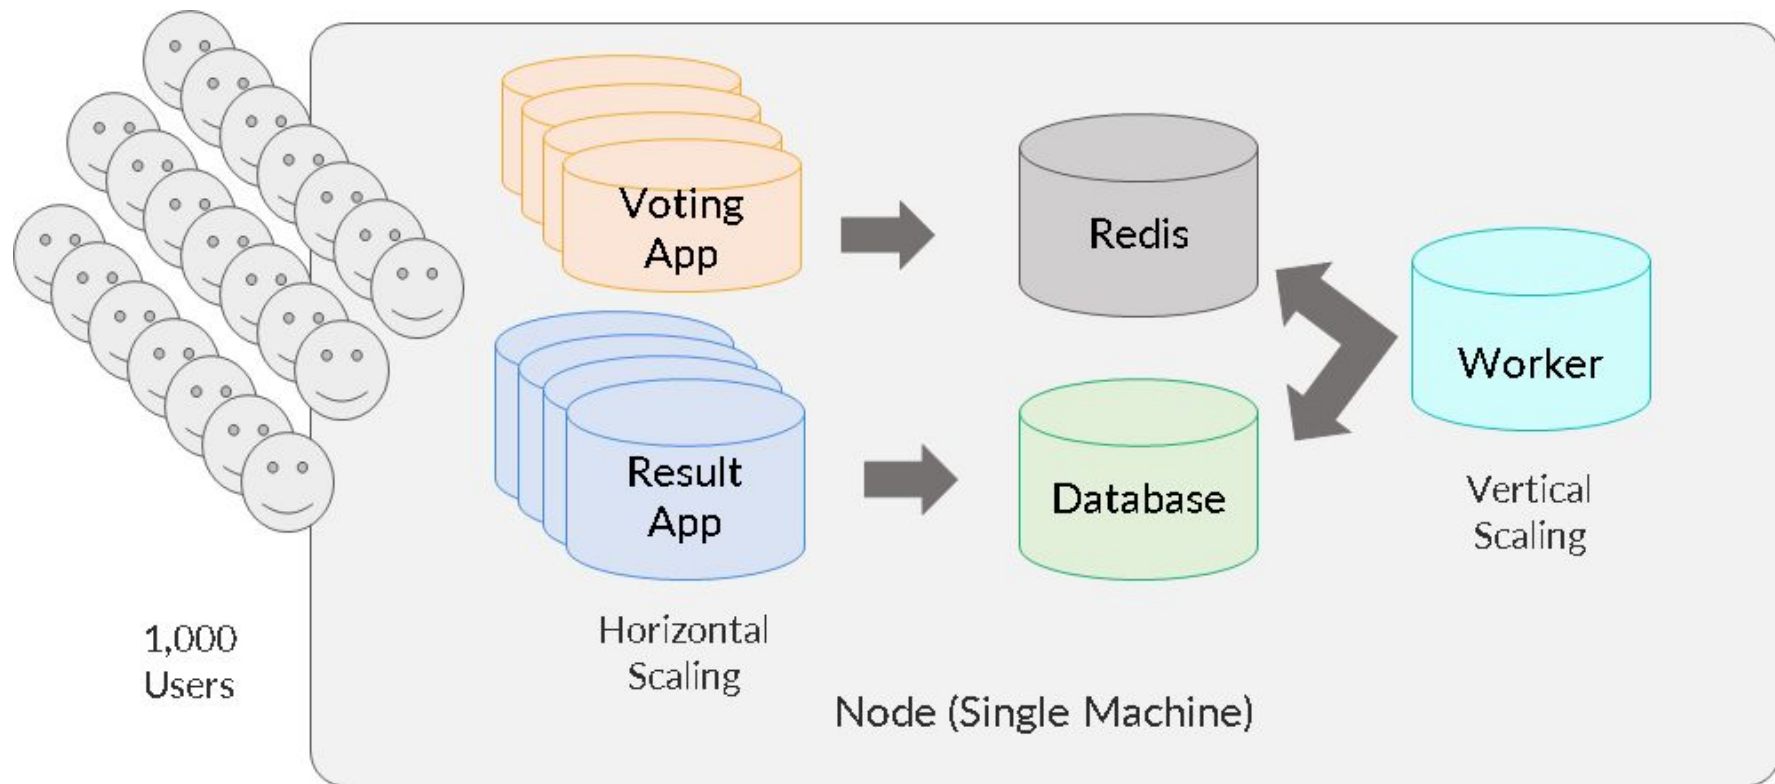
○ Horizontal

- No downtime. Dynamically increase or decrease the resources.
- You can choose to run only specific number of hardwares.
- Easier to run fault-tolerance.
- Infinite scope of scaling.
- Architectural design can become highly complicated.

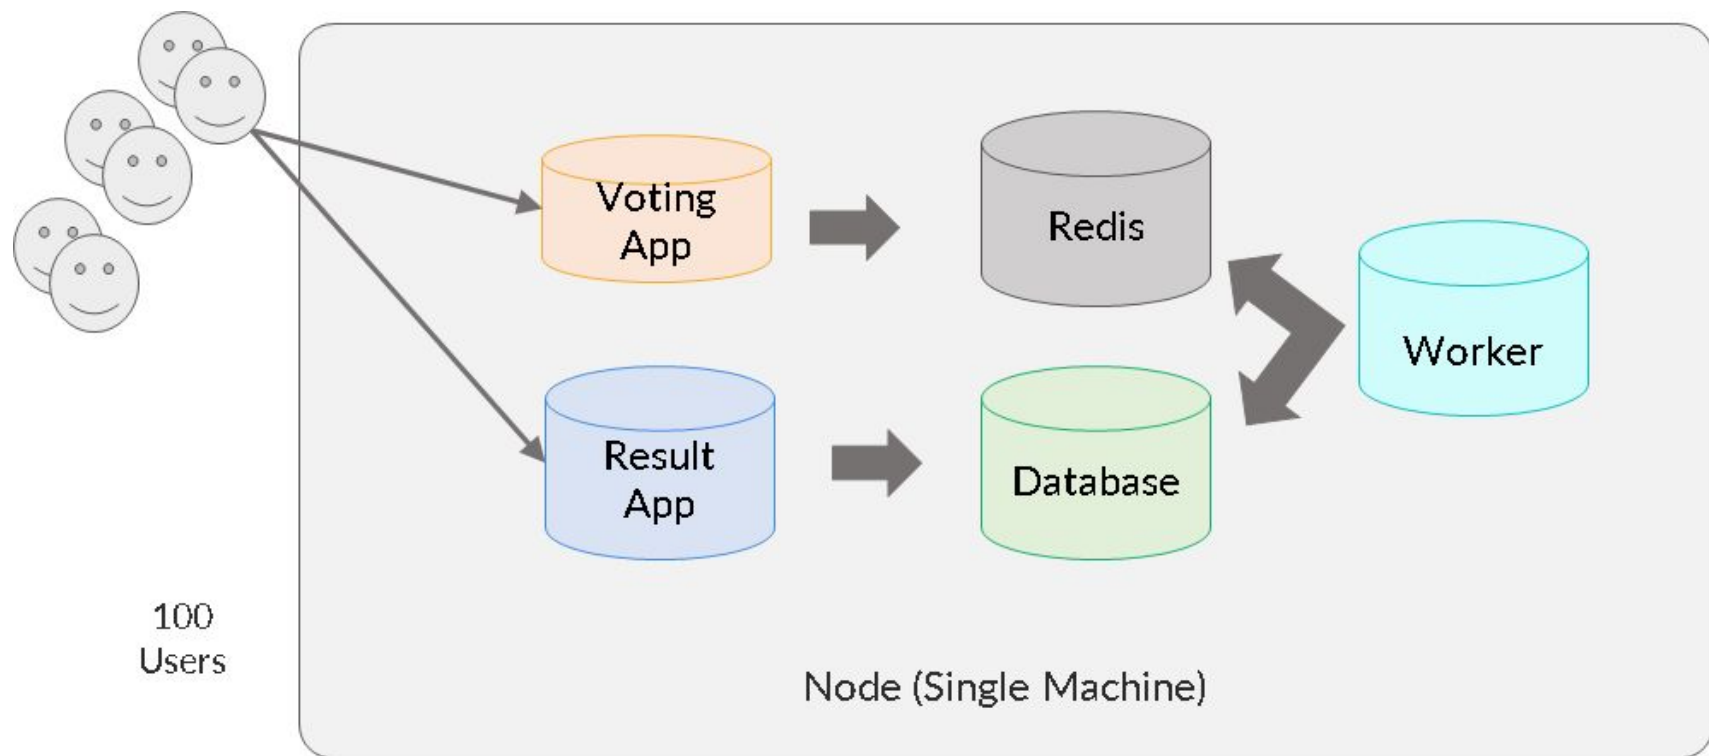
DEMO 2 : Scaling

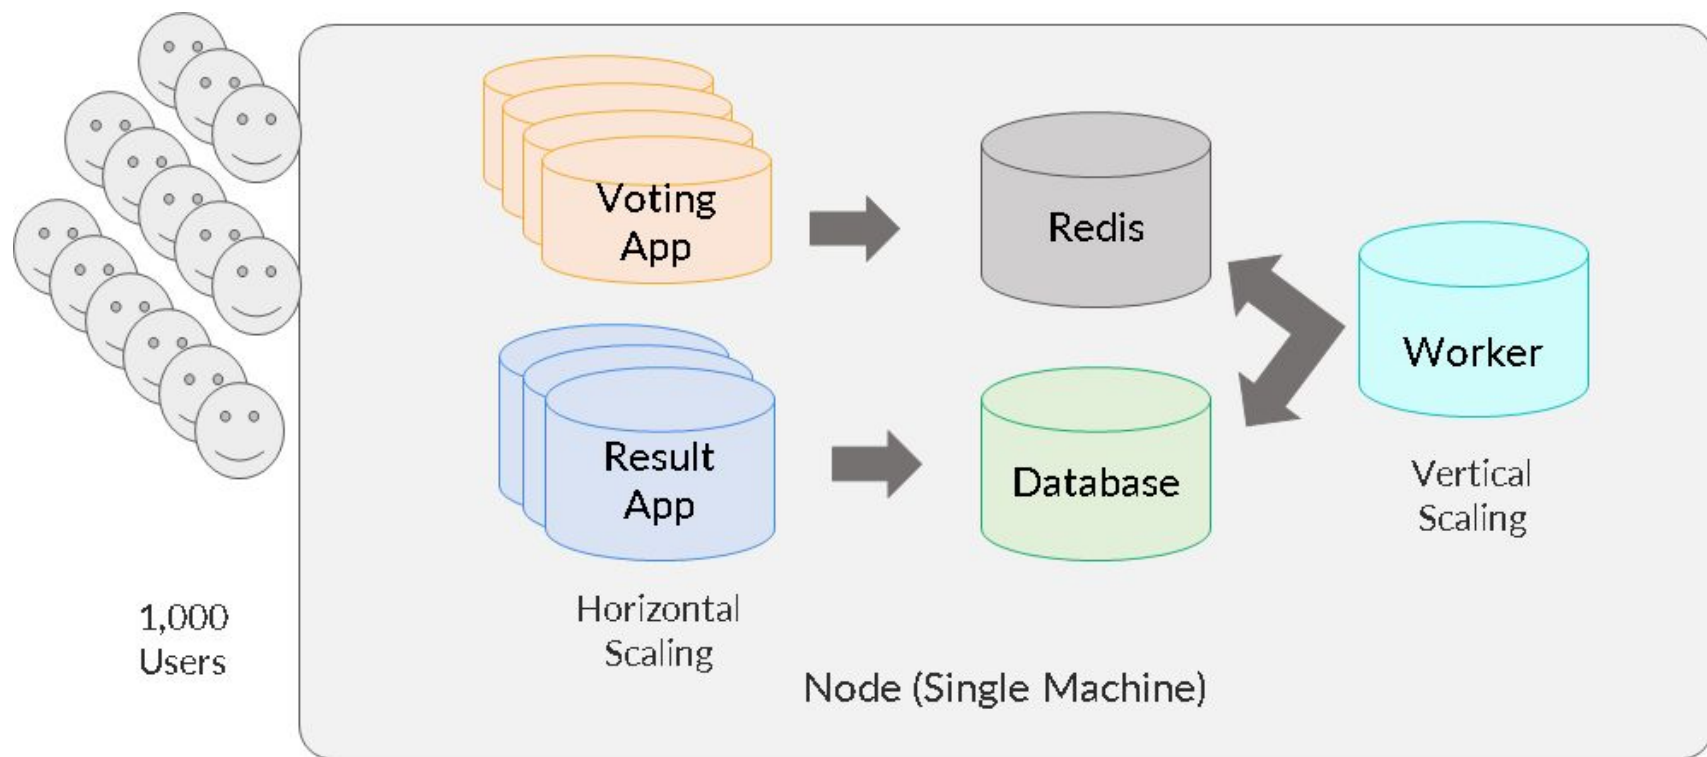
- Let's see how we can scale application in AWS environment
- Vertical scaling in EC2 & RDS
- Horizontal scaling using ASG (auto scaling groups)

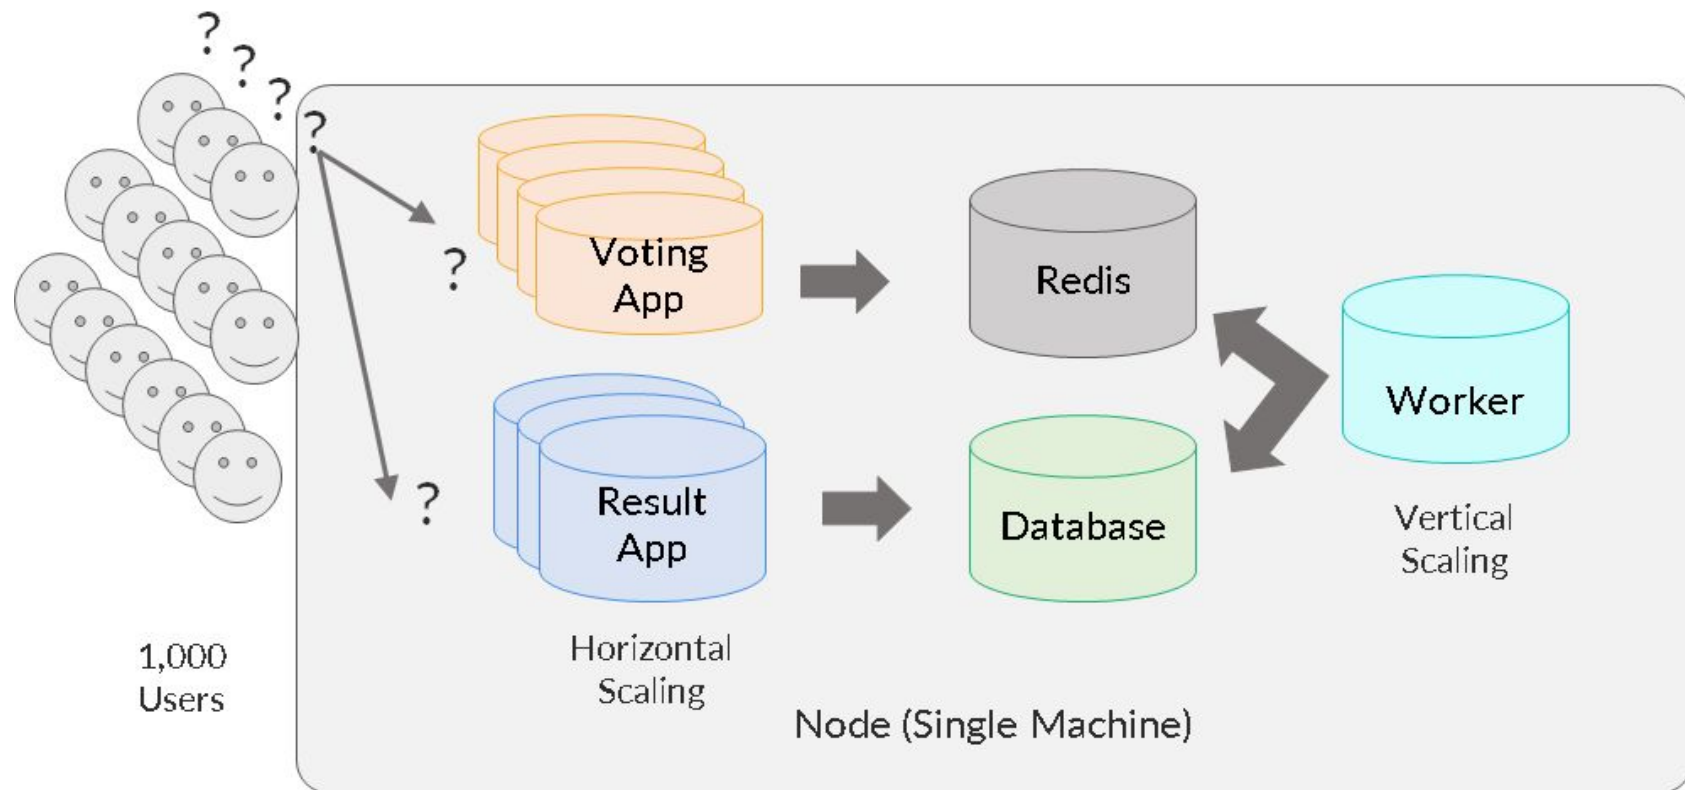
Scaling - What do We Achieve?



- Manual
 - Requires manual efforts for any scaling operation
 - Not ideal for sudden fluctuations
 - Can lead to human errors
- Scheduled
 - For predictable workloads
 - Effective for non-production system or predictable production traffic
 - Example: Shut down of non production environment in night hours
 - Cost-effective
- Automatic
 - Based on predefined rules and thresholds
 - Makes an application always available
 - Must be configured in dynamic traffic applications

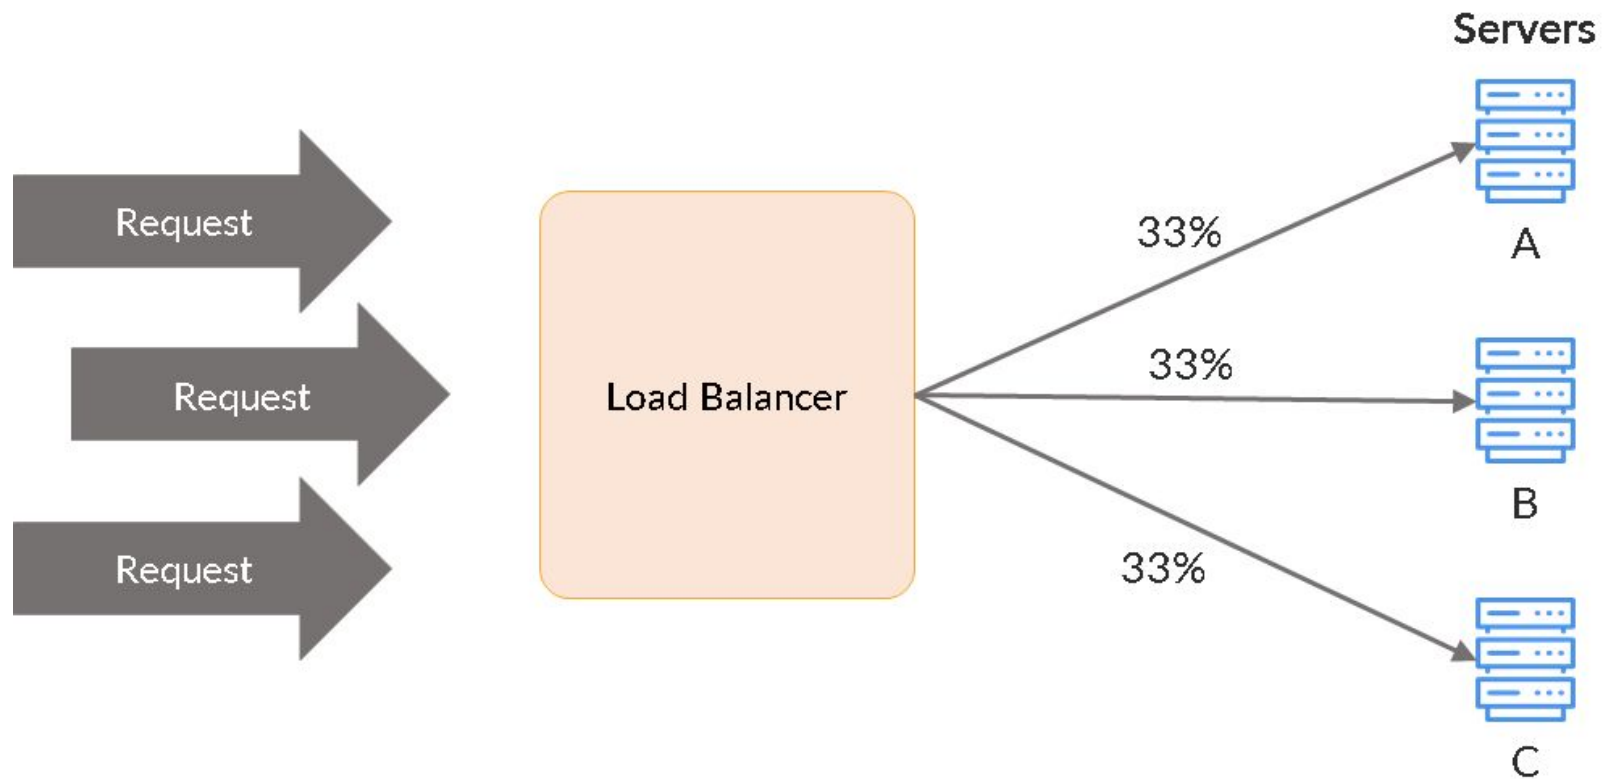






- Load balancing refers to efficiently distributing application requests across all the healthy back-end servers.
- It is the key component of highly available infrastructure.
- It provides the flexibility to add or remove servers as per the demand.
- The user accesses the LB, which forwards the user's request to a back-end server.
- It saves us from launching any application server in a public subnet. Only the LB needs to be launched in a public subnet.
- It ensures effective use of all the available servers.
- It saves cost.

Load Balancing (LB)

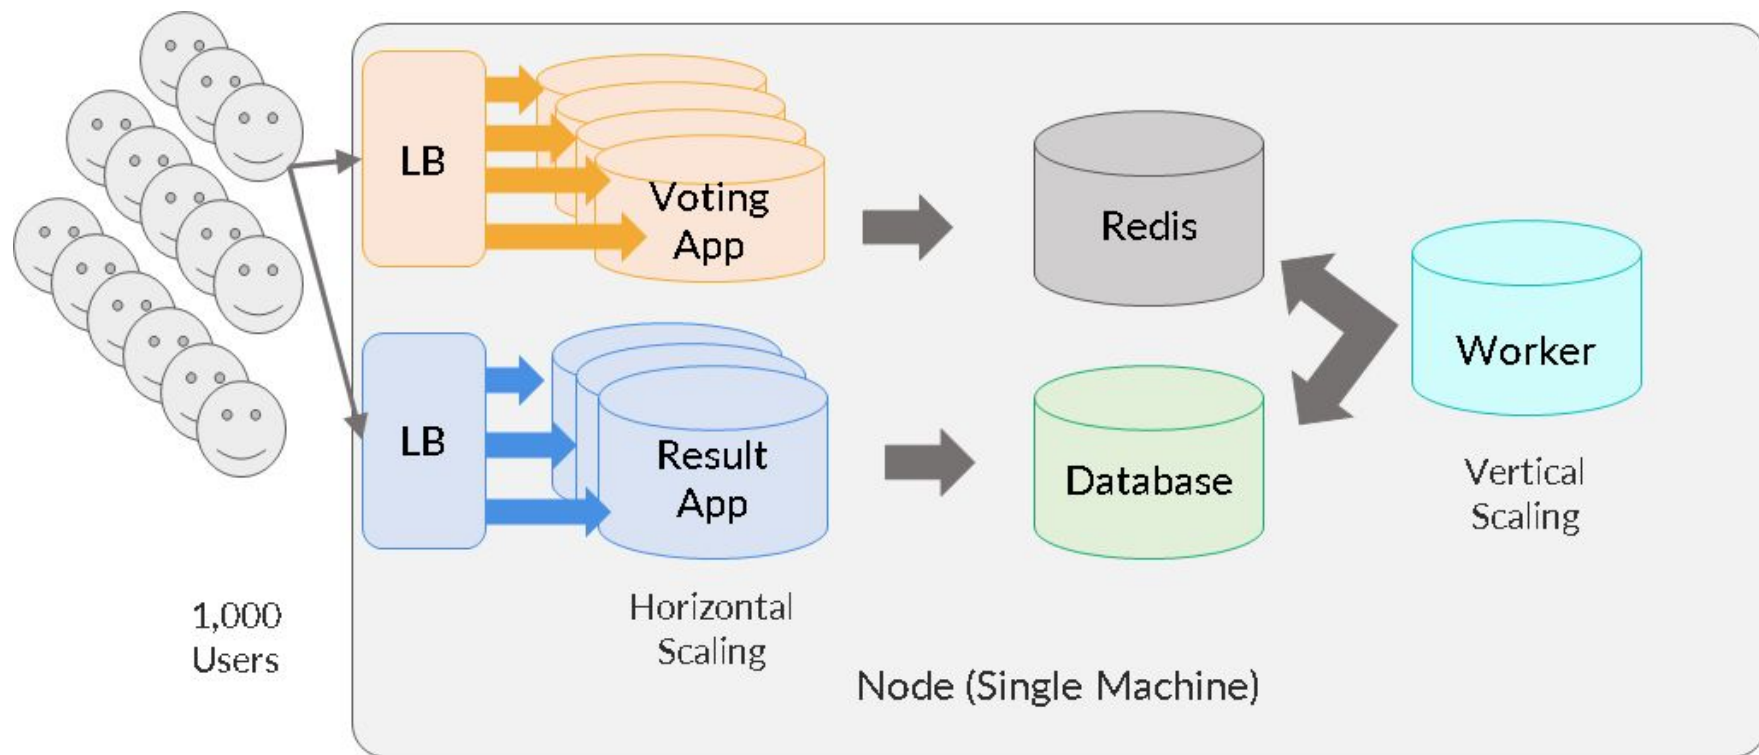


- **Round robin** - This is the most common type and is widely used. Requests are distributed sequentially [one by one to each server].
- **Least connections** - The server with the fewest connections is awarded the new request.
- **Least time** - It depends on the fastest response time and the fewest active connections.
- **IP hash** - This is used for sticky sessions. The IP address of the client is used to determine which server receives the request.

DEMO 3 : Load Balancing

- Use of Nginx for load balancing
- Implementing load balancing types using Nginx
- Load balancing in AWS environment

Load Balancing - What do We Achieve?



Poll 3 (15 seconds)

Rahul is handling an internal attendance marking application on an EC2 machine for his company. All the employees work in the 10 am–7 pm IST shift. From the options below, Rahul wants to select the one that would result in the minimum cost.

- A. Rely upon a manual scaling system
- B. Go for scheduled scaling with a pre-decided machine size
- C. Opt for automatic scaling when CPU or memory reaches 60%

Poll 3 (15 seconds)

Rahul is handling an internal attendance marking application on an EC2 machine for his company. All the employees work in the 10 am–7 pm IST shift. From the options below, Rahul wants to select the one that would result in the minimum cost.

- A. Rely upon a manual scaling system
- B. Go for scheduled scaling with a pre-decided machine size**
- C. Opt for automatic scaling when CPU or memory reaches 60%

Poll 4 (15 seconds)

Jack sets up five back-end services with Nginx for load balancing. Many users in different locations reported latency related issues. Jack checked the status of the system and found that all the systems were underutilised. What he can try in order to resolve the issue?

- A. Add two more machines to the cluster
- B. Use the IP hash approach for load balancing
- C. Use the least connections approach for load balancing
- D. Ask the users to check their Internet bandwidth

Poll 4 (15 seconds)

Jack sets up five back-end services with Nginx for load balancing. Many users in different locations reported latency related issues. Jack checked the status of the system and found that all the systems were underutilised. What he can try in order to resolve the issue?

- A. Add two more machines to the cluster
- B. Use the IP hash approach for load balancing
- C. Use the least connections approach for load balancing**
- D. Ask the users to check their Internet bandwidth

Advanced Features of Docker Swarm

- Prevents a node from receiving new tasks
- Will stop existing tasks running on that node and launch replicas of those tasks on other nodes
- Helps with patching servers and debugging issues

```
docker node update --availability drain worker1
```

```
docker node update --availability active worker1
```

- `docker service create --name webserver --replicas 5 nginx`
- `docker service ps webserver`
 - See where these 5 tasks have been placed
- Now assume we want to take instance 172.31.1.160 in maintenance mode

```
ubuntu@ip-172-31-5-46:~$ docker service ps webserver
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
drzgpqngahas	webserver.1	nginx:latest	ip-172-31-5-46	Running	Running 29 seconds ago		
vlt10u00zknn	webserver.2	nginx:latest	ip-172-31-1-160	Running	Running 28 seconds ago		
lbd4iiovszio	webserver.3	nginx:latest	ip-172-31-5-46	Running	Running 29 seconds ago		
p4ei5mohmkq2	webserver.4	nginx:latest	ip-172-31-1-160	Running	Running 28 seconds ago		
c2c2n8kkppkb	webserver.5	nginx:latest	ip-172-31-1-160	Running	Running 28 seconds ago		

```
ubuntu@ip-172-31-5-46:~$
```


- `docker node update --availability drain ip-172-31-1-160`
- `docker service ps webserver`
 - Now, See where these 5 tasks have been placed

```
ubuntu@ip-172-31-5-46:~$
```

```
ubuntu@ip-172-31-5-46:~$ docker service ps webserver
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
drzgpqngahas	webserver.1	nginx:latest	ip-172-31-5-46	Running	Running 5 minutes ago		
qu0ym1eajlnw	webserver.2	nginx:latest	ip-172-31-5-46	Running	Running about a minute ago		
vlt10u00zknn	_ webserver.2	nginx:latest	ip-172-31-1-160	Shutdown	Shutdown about a minute ago		
lbd4iiovszio	webserver.3	nginx:latest	ip-172-31-5-46	Running	Running 5 minutes ago		
guheqtq25j6i	webserver.4	nginx:latest	ip-172-31-5-46	Running	Running about a minute ago		
p4ei5mohmkq2	_ webserver.4	nginx:latest	ip-172-31-1-160	Shutdown	Shutdown about a minute ago		
unmh49pcu2lx	webserver.5	nginx:latest	ip-172-31-5-46	Running	Running about a minute ago		
c2c2n8kkppkb	_ webserver.5	nginx:latest	ip-172-31-1-160	Shutdown	Shutdown about a minute ago		

```
ubuntu@ip-172-31-5-46:~$
```


- “**Inspect**” provides detailed information about any service or node.
- By default, it will show the results in a JSON array.
- It helps with debugging an issue (ex- container not getting scheduled or a node not getting registered.)

❏ `docker service inspect serviceName`

❏ **`docker service inspect --pretty serviceName`**
(display the details about a service in an easily readable format.)

❏ `docker node inspect <nodeId> --pretty`

- The Swarm service provide a few different ways for you to control scale and placement of services on different nodes:
 - **Replicated** and **global** modes of services
 - **Resource constraints** like requirement of CPU/memory
 - **Placement constraints** (ex- Run container only on a node with label upgrad=true)
 - **Placement preferences** (ex- spread containers evenly across groups of nodes where grouping is based on a label's value)

- Let's place a service using placement constraint
 - Constraint : Node should have label => region==india
 - docker service create --name webserver --constraint node.labels.region==india --replicas 3 nginx
 - If no node fulfil this constraint, the containers will go into pending state
 - See the current status using, docker service ps webserver

```
ubuntu@ip-172-31-5-46:~$
```

```
ubuntu@ip-172-31-5-46:~$ docker service ps webserver
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
vfu5kb1frqro	webserver.1	nginx:latest		Running	Pending 3 minutes ago	"no suitable node (scheduling ..."	
v4bxd3i3lsdr	webserver.2	nginx:latest		Running	Pending 3 minutes ago	"no suitable node (scheduling ..."	
8joutz9m9i0q	webserver.3	nginx:latest		Running	Pending 3 minutes ago	"no suitable node (scheduling ..."	

```
ubuntu@ip-172-31-5-46:~$
```

```
ubuntu@ip-172-31-5-46:~$
```

- Set node label to any one node using
 - `docker node update --label-add region=india <NODE ID>`
- Now, see current status of containers
- All containers will be on same single node to which we applied the label

```
ubuntu@ip-172-31-5-46:~$
```

```
ubuntu@ip-172-31-5-46:~$ docker node update --label-add region=india ip-172-31-1-160
ip-172-31-1-160
```

```
ubuntu@ip-172-31-5-46:~$ docker service ps webserver
```

ID	NAME	IMAGE	NODE	DESIRED STATE	CURRENT STATE	ERROR	PORTS
vfu5kb1frqro	webserver.1	nginx:latest	ip-172-31-1-160	Running	Running 1 second ago		
v4bxd3i3lsdr	webserver.2	nginx:latest	ip-172-31-1-160	Running	Running 1 second ago		
8joutz9m9i0q	webserver.3	nginx:latest	ip-172-31-1-160	Running	Running 1 second ago		

```
ubuntu@ip-172-31-5-46:~$
```

```
ubuntu@ip-172-31-5-46:~$
```

```
ubuntu@ip-172-31-5-46:~$ █
```

- Let us set up the voting application with all the feature implementations that we have discussed till now:
 - Set up voting application on a three-node cluster
 - Scale the services
 - Apply node constraints on the key value and the resource basis
 - Set up Nginx on the server

Demo 4 :

Docker Swarm Advanced Features

- Node draining
- Placement constraints and placement preferences
- Load balancing in AWS environment

Demo 5 :

Docker Voting Application

Important Concepts and Questions

1. What are the differences between a forward proxy and a reverse proxy?
2. How can you host multiple websites using Nginx?
3. What is scaling? Explain the different strategies and the types of scaling.
4. Explain the different load balancing techniques. Write a practical use case of each.
5. What is the use of draining?
6. Can you explain the different placement constraints present in Docker Swarm?

Doubt Clearance Window

Session 1.1 covered these topics:

1. Docker basic concepts recap
2. Setting up an application using Docker Compose
3. Issues with Docker Compose
4. What and why: Docker orchestration
5. Introduction to Docker Swarm and its architecture
6. Deploying an application in a Swarm cluster
7. Basic features of Swarm
8. Docker Compose versus Swarm

This class covered the following topics:

1. Nginx: Reverse proxy versus proxy, Nginx folder structure, virtual site hosting
2. Scaling and load balancing
3. Understand why, when and how to scale an application
4. Horizontal versus vertical scaling
5. Load balancing types
6. Schedule services based on different strategies and constraints
7. Deploy an application on a Swarm cluster

What's Next??

1. What is ECS, ECS Components
2. Concepts of ECR, EC2 and Fargate; task definition; and services
3. ECS pricing
4. Launch Getting started Cluster, understand things, update services and task definition
5. Launch your own Fargate cluster. Create your own TD and services
6. Logging and monitoring
7. Update services
8. ECS load balancing with ALB
9. Deploy containers in ECS Fargate mode

1. Create a docker file such that when a docker image is build using this file, and a container is run using that image, it should run the Nginx web server inside it and replace the welcome page text with container IP. Now launch three replicas of the same container by docker run command. Launch a new Nginx container and change its configuration file such that it proxy pass the traffic to the other three containers. Try different load balancing types in the configuration file.
2. Host a site yourname.com. Create an index page and a fail-safe page. If a page corresponding to the URI is not available, then the fail-safe page is served.
3. Set up the voting application using Docker Swarm:
 - a. Set up the application on a three-node cluster
 - b. Scale the voting and result services to 5 replicas
 - c. [Apply node constraints on the key value and the resource basis](#)
 - d. Now add Nginx container in the stack so that voting and result frontend can be accessed using votingapp.com & resultapp.com [Hint: you also have to add two DNS entries in /etc/hosts file in your local system mapped to ec2 IP.
4. [Read about the difference between a placement constraint and a placement preference.](#)

Tasks to Complete After Today's Session

MCQs
Practice Questions
Practice Project



Thank You!